

# DEPLOYMENT STRATEGIES FOR YOUR MICRO-SERVICE WITH K8S



+ HOW TO PREPARE YOUR SERVICE

Wojciech Barczynski - SMACC.io | Hypatos.ai  
Listopad 2018

# WOJCIECH BARCZYŃSKI

- Lead Software Engineer  
& System Engineer
- Interests:  
working software
- Hobby:  
teaching software  
engineering

**SMACC**

# BACKGROUND

- ML FinTech ➔ micro-services and k8s
- Before:
  - 1 z 10 Indonesian mobile e-commerce (Rocket Internet)
- Spent 3.5y with Openstack, 1000+ nodes, 21 data centers
- I do not like INFRA :D

# STORY

- Lyke - [12.2016 - 07.2017]
- SMACC - [10.2017 - present]

# AGENDA.

- Mikroserwises
- Dlaczego kubernetes?
- Strategie deploymentu
- Jak przygotować mikroserwis?

# MICRO-SERVICES

- Scalling the team
- ...
- Scalling the products

# MICRO-SERVICES

- Independent
- Following 12factor app
- Self-aware

# 12FACTOR APP

- Heroku - 2011
- App easy to run in production
- Low TCO
- Easy to manage

# MICRO-SERVICES

- */healthz*
- */metrics*
- */readiness*
- */info*

# MICRO-SERVICES

- Do not need to share code
- Share the same conventions
- Every git repos looks familiar
- Might be in different tech

# MICRO-SERVICES

- Not a silver bullet :)

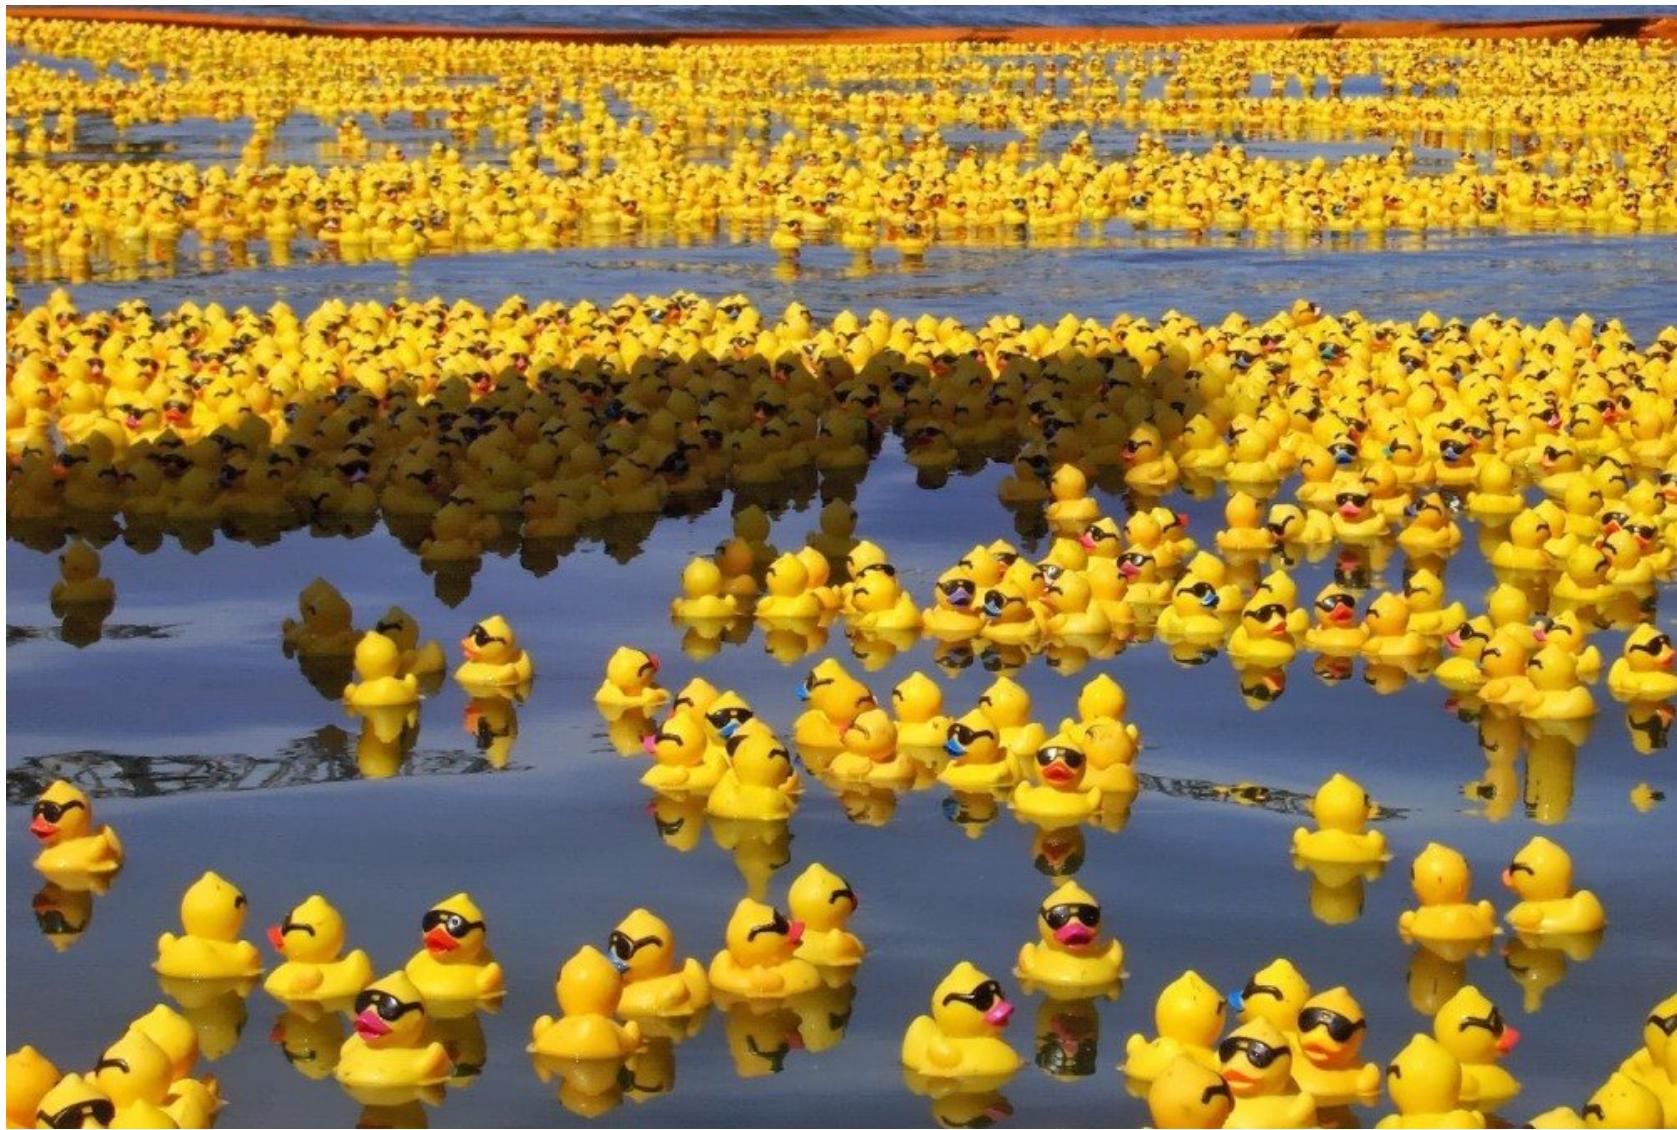
# KUBERNETES



# WHY?

- Administracja jest trudna i kosztowna
- Virtualne Maszyny, ansible, salt, etc.
- Za dużo ruchomych części
- Nie kończąca się standaryzacja

# MIKROSERWISY AAA!



# WHY?

- Cloud is not so cheap - \$\$\$

# IMAGINE

- do not need to think about IaaS
- no login on a VM
- less gold plating your CI / CD ...
- DC as a black box

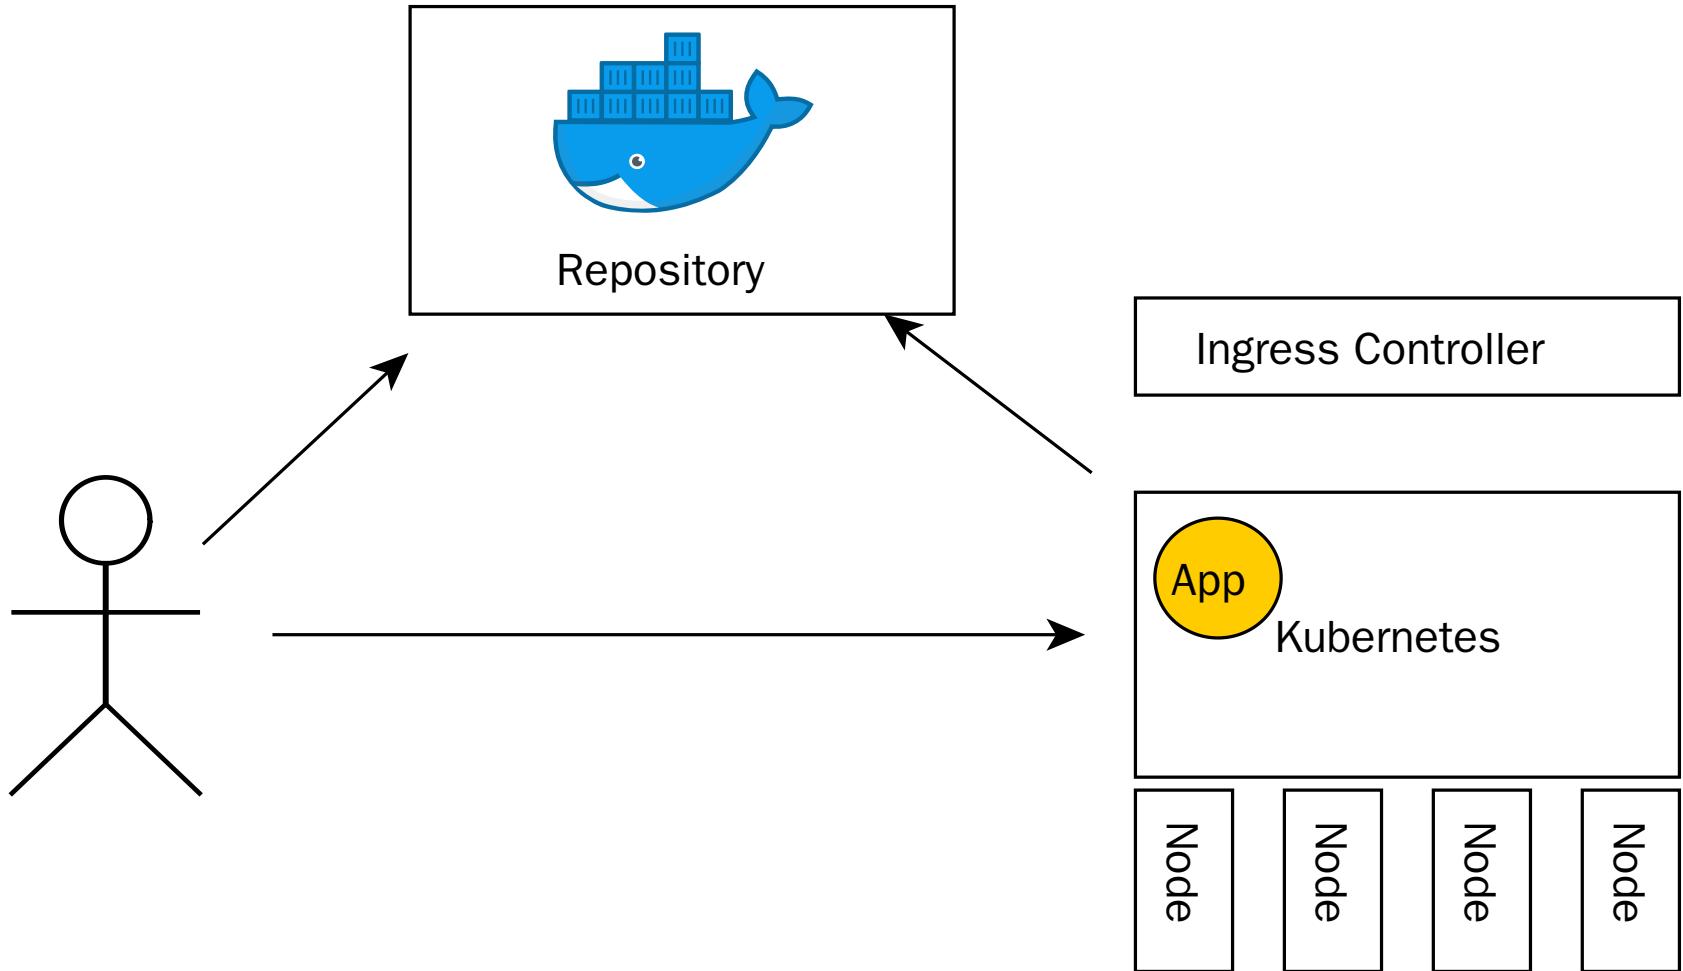
# KUBERNETES

- Container management
- Service and application mindset
- Simple Semantic\*
- Independent from IaaS provider

## KUBERNETES

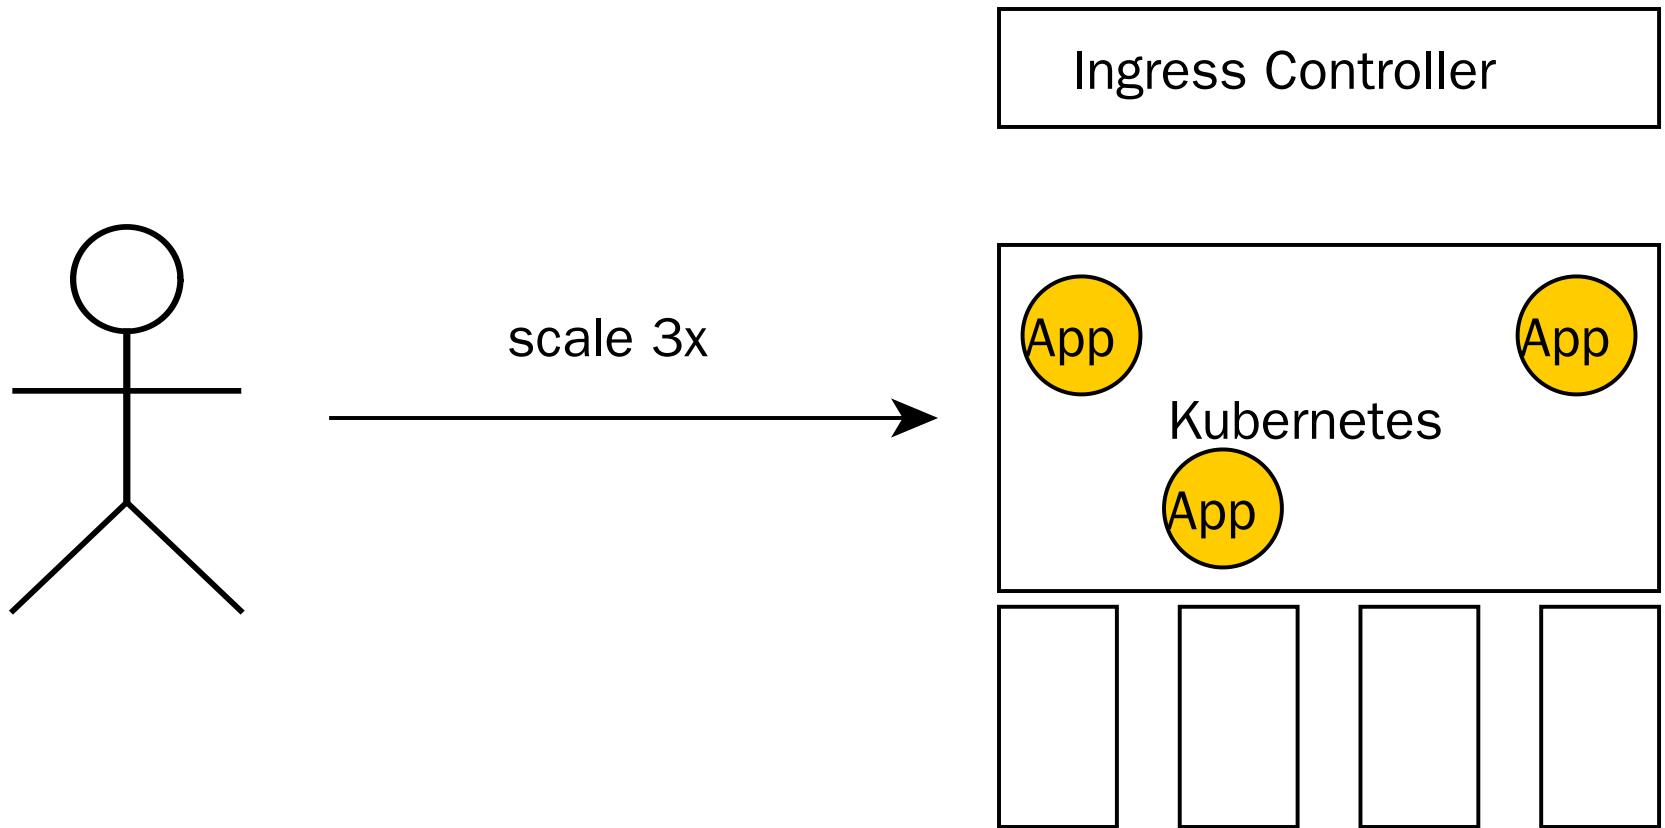
- Batteries for your 12factory apps
- Service discovery, meta-data support
- Utilize resources to nearly 100%

# KUBERNETES



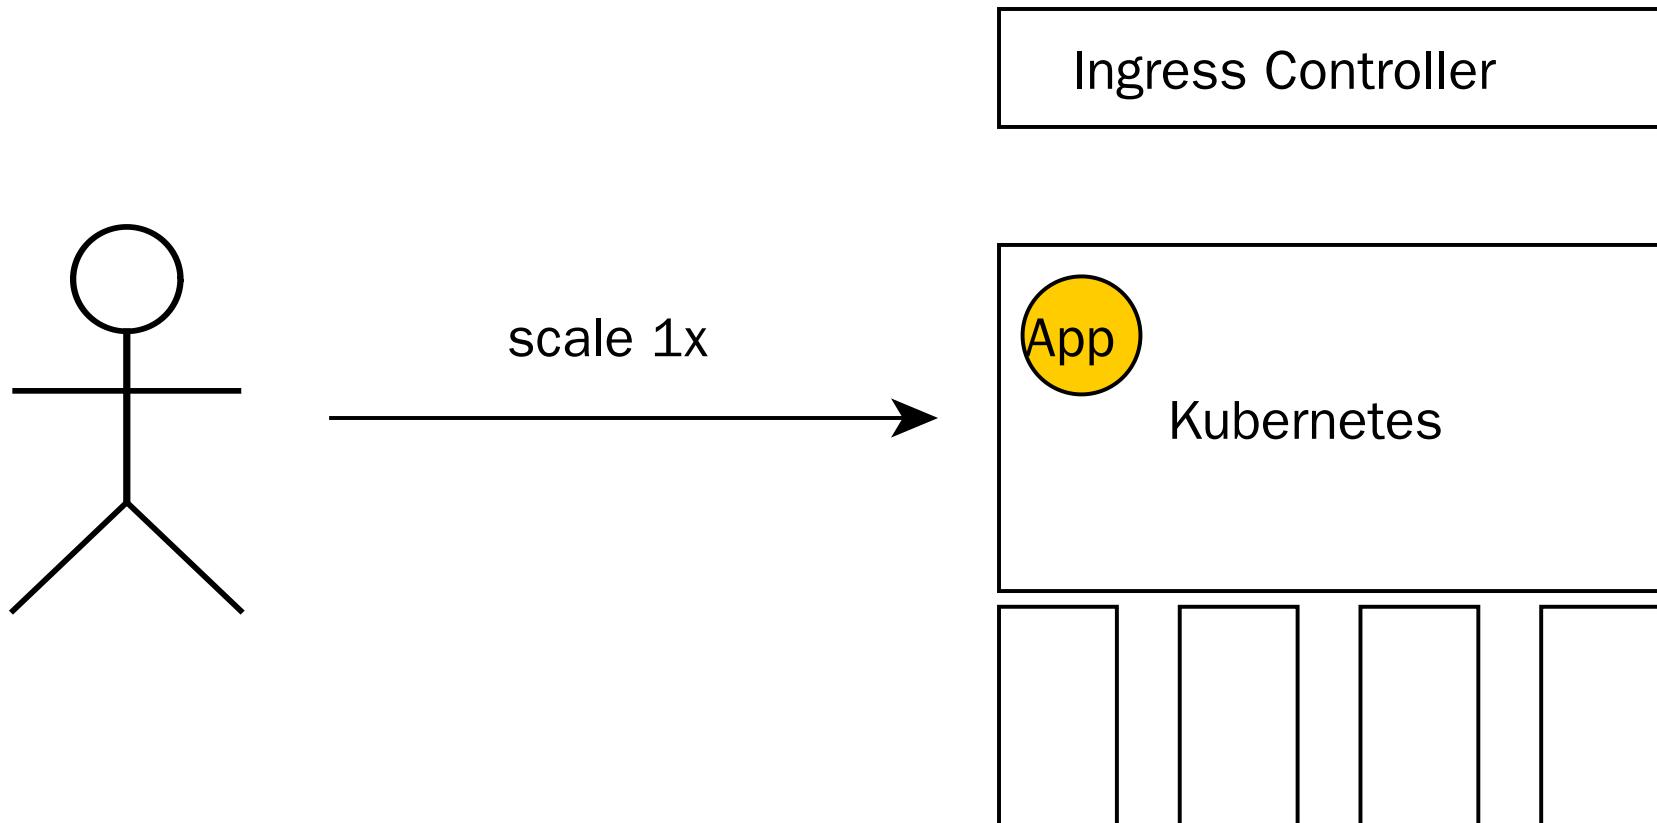
make docker\_push; kubectl create -f app-srv-dpl.yaml

# SCALE UP! SCALE DOWN!



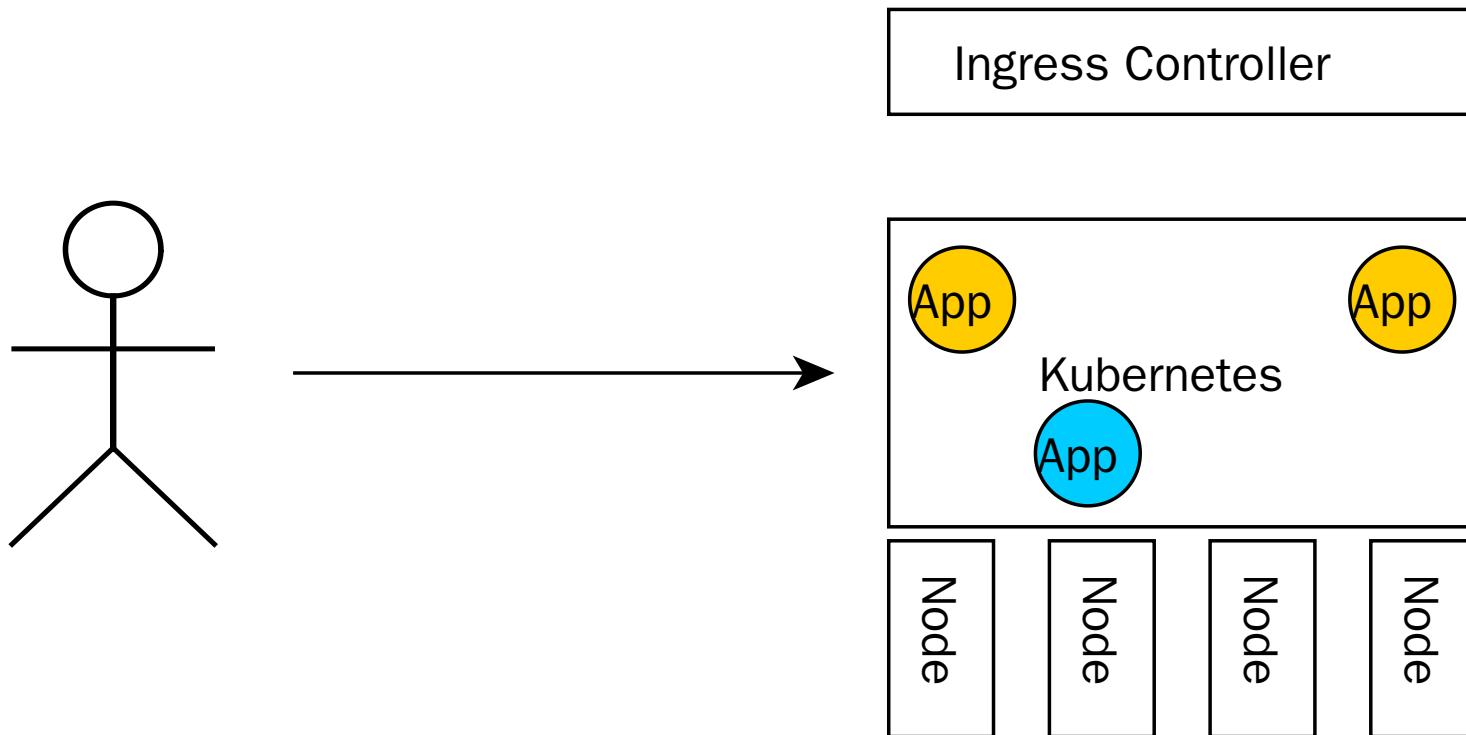
```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

# SCALE UP! SCALE DOWN!



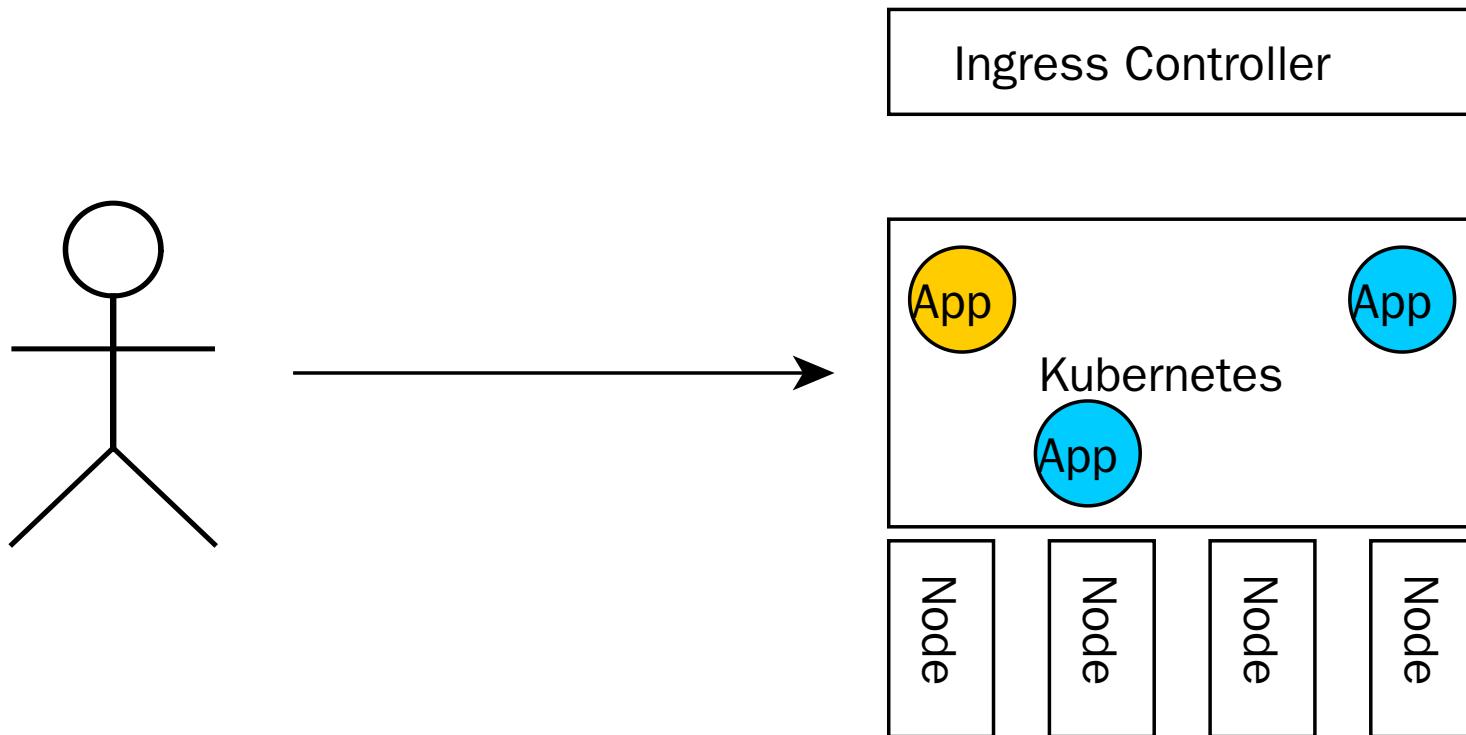
```
kubectl --replicas=1 -f app-srv-dpl.yaml
```

# ROLLING UPDATES!



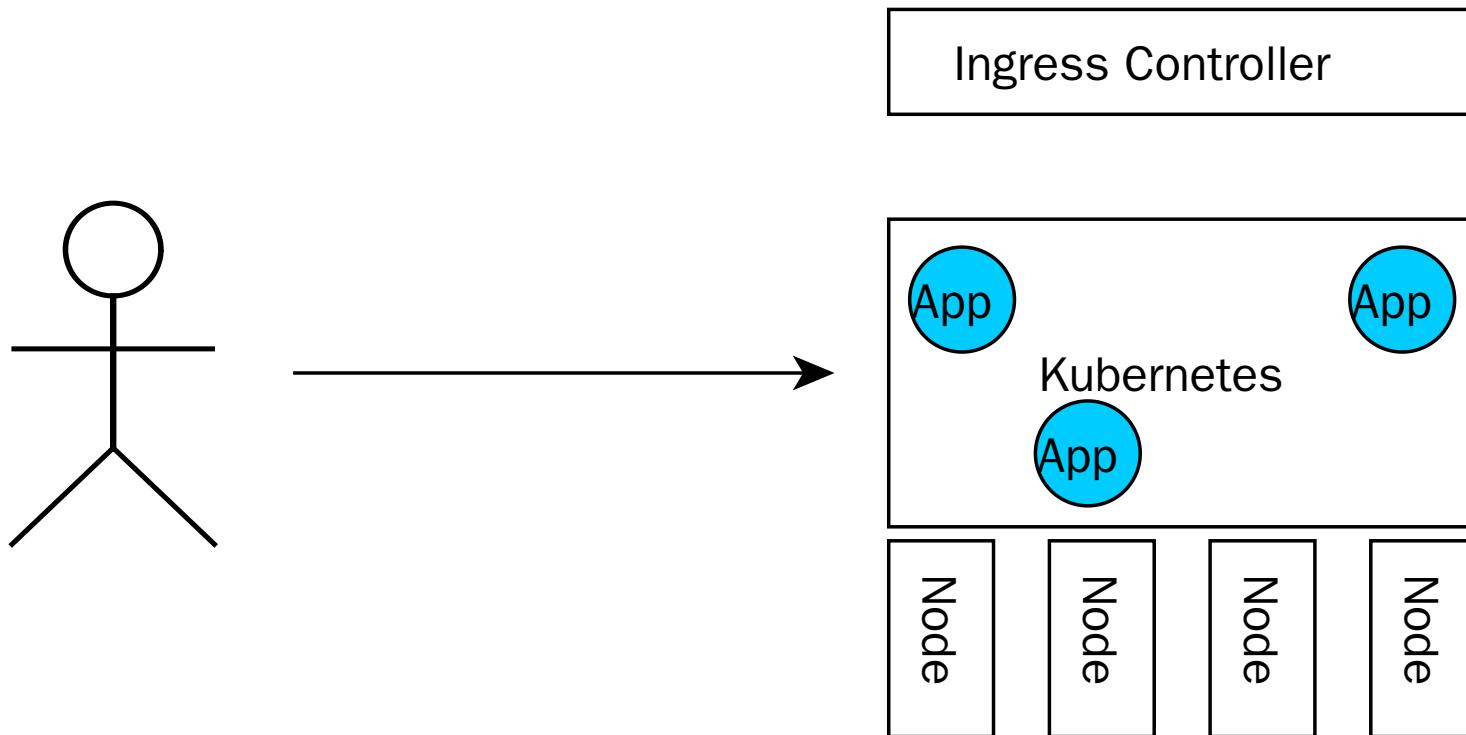
```
kubectl set image deployment/app app=app:v2.0.0
```

# ROLLING UPDATES!



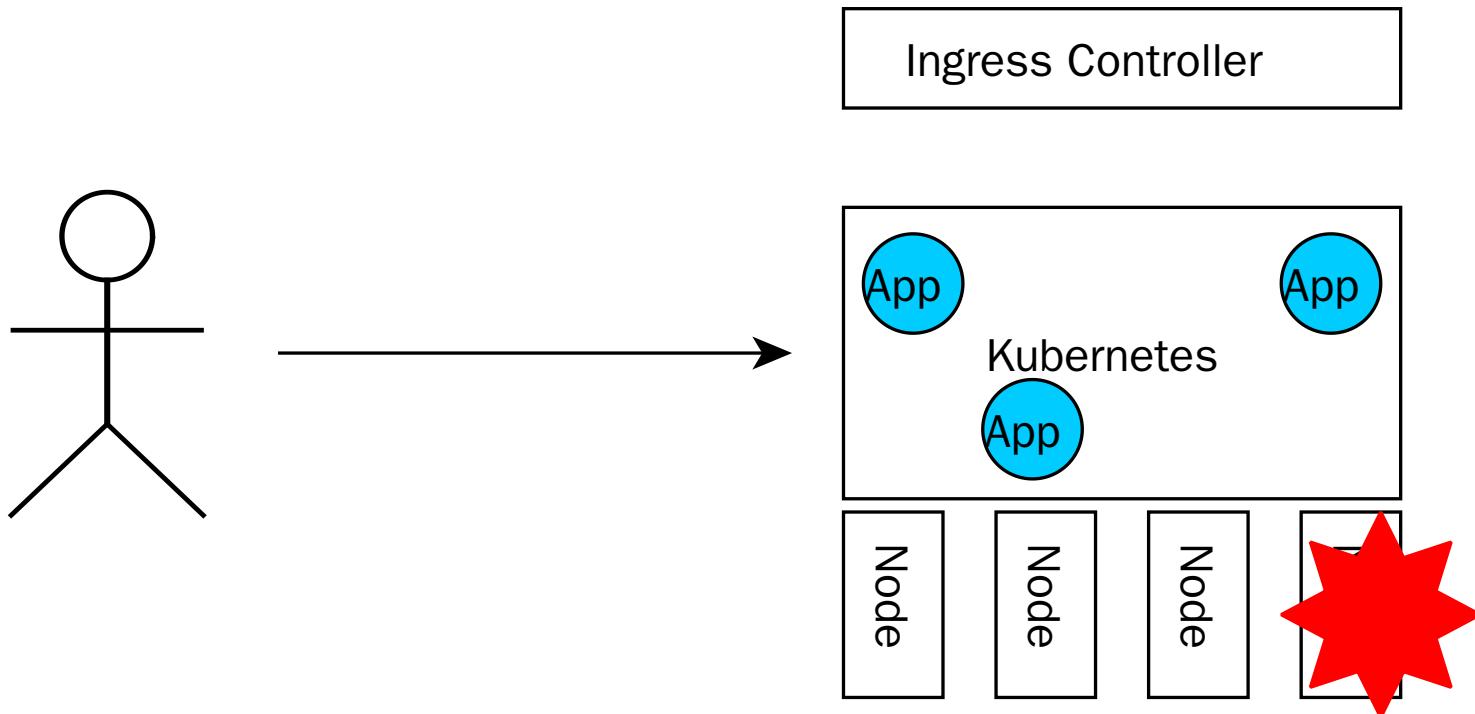
```
kubectl set image deployment/app app=app:v2.0.0
```

# ROLLING UPDATES!

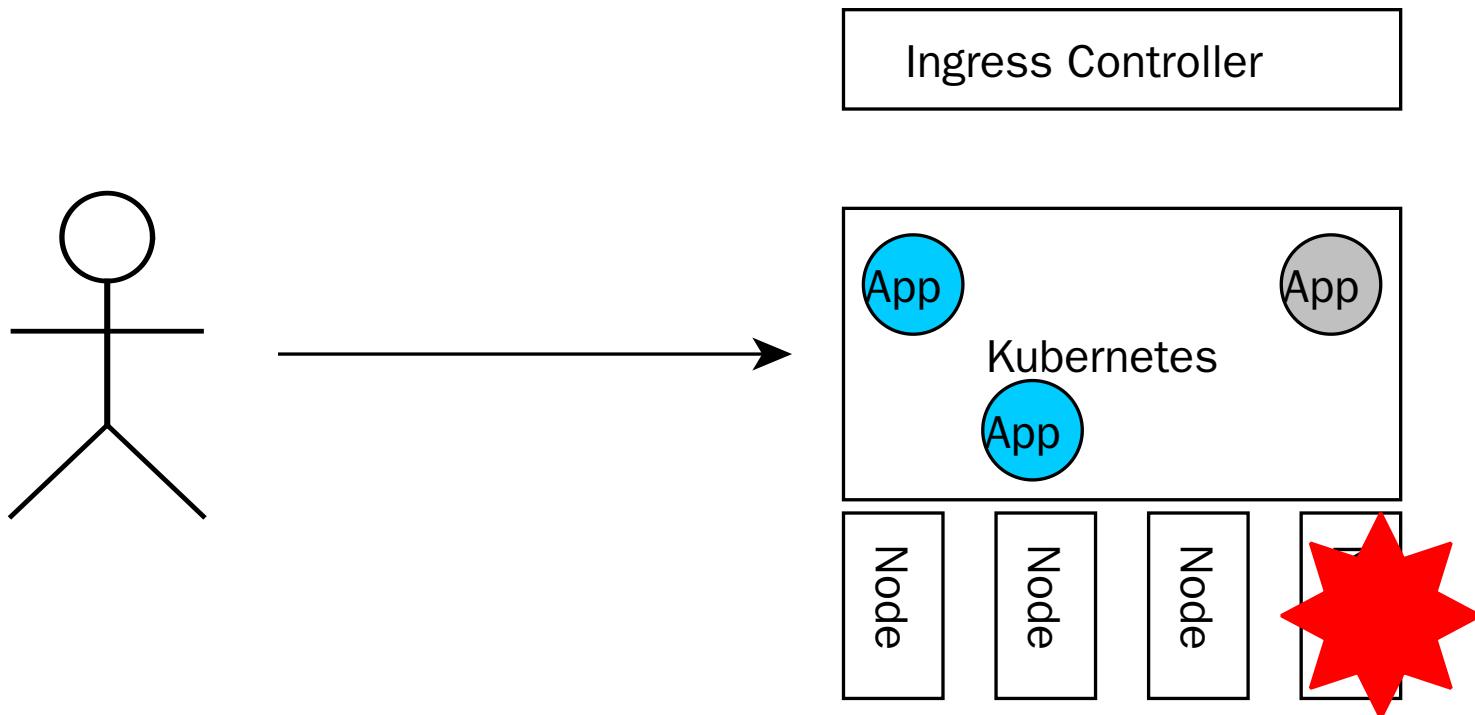


```
kubectl set image deployment/app app=app:v2.0.0
```

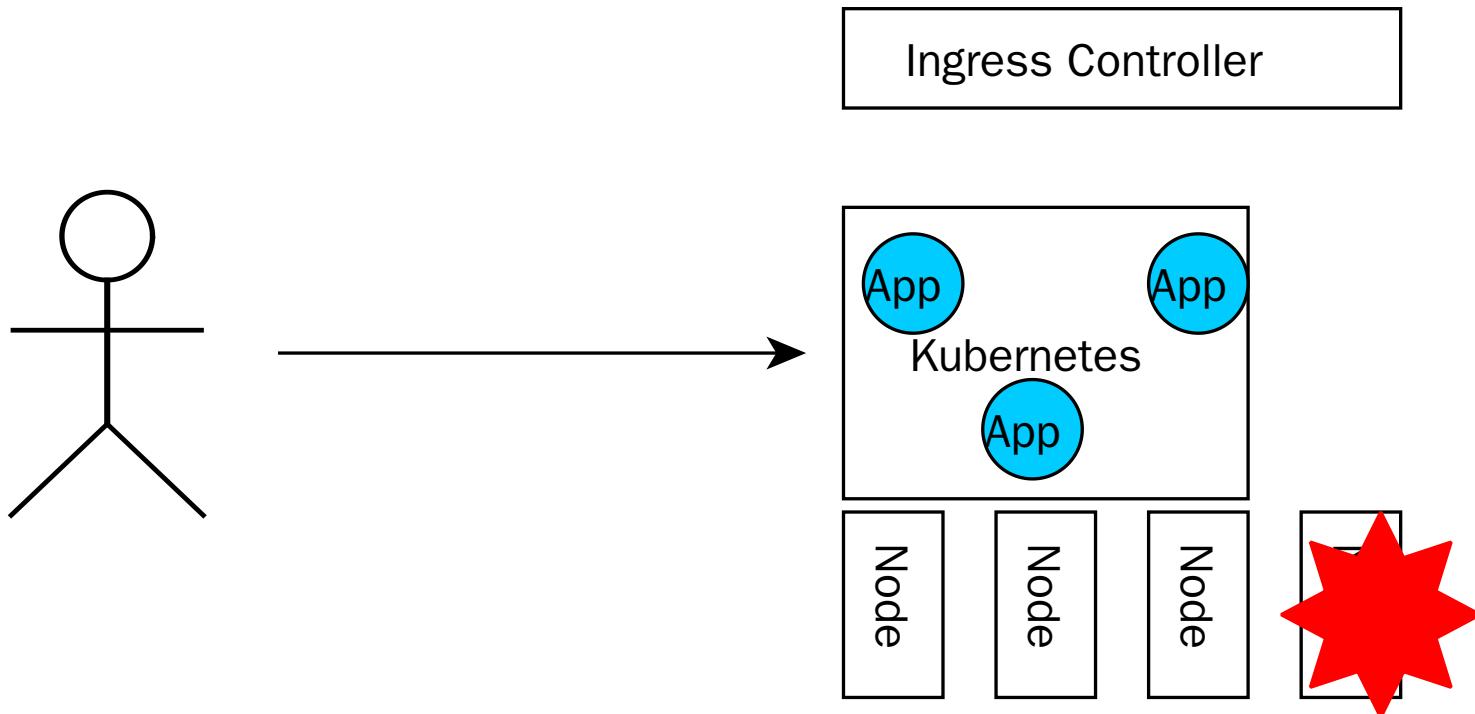
# RESISTANCE!



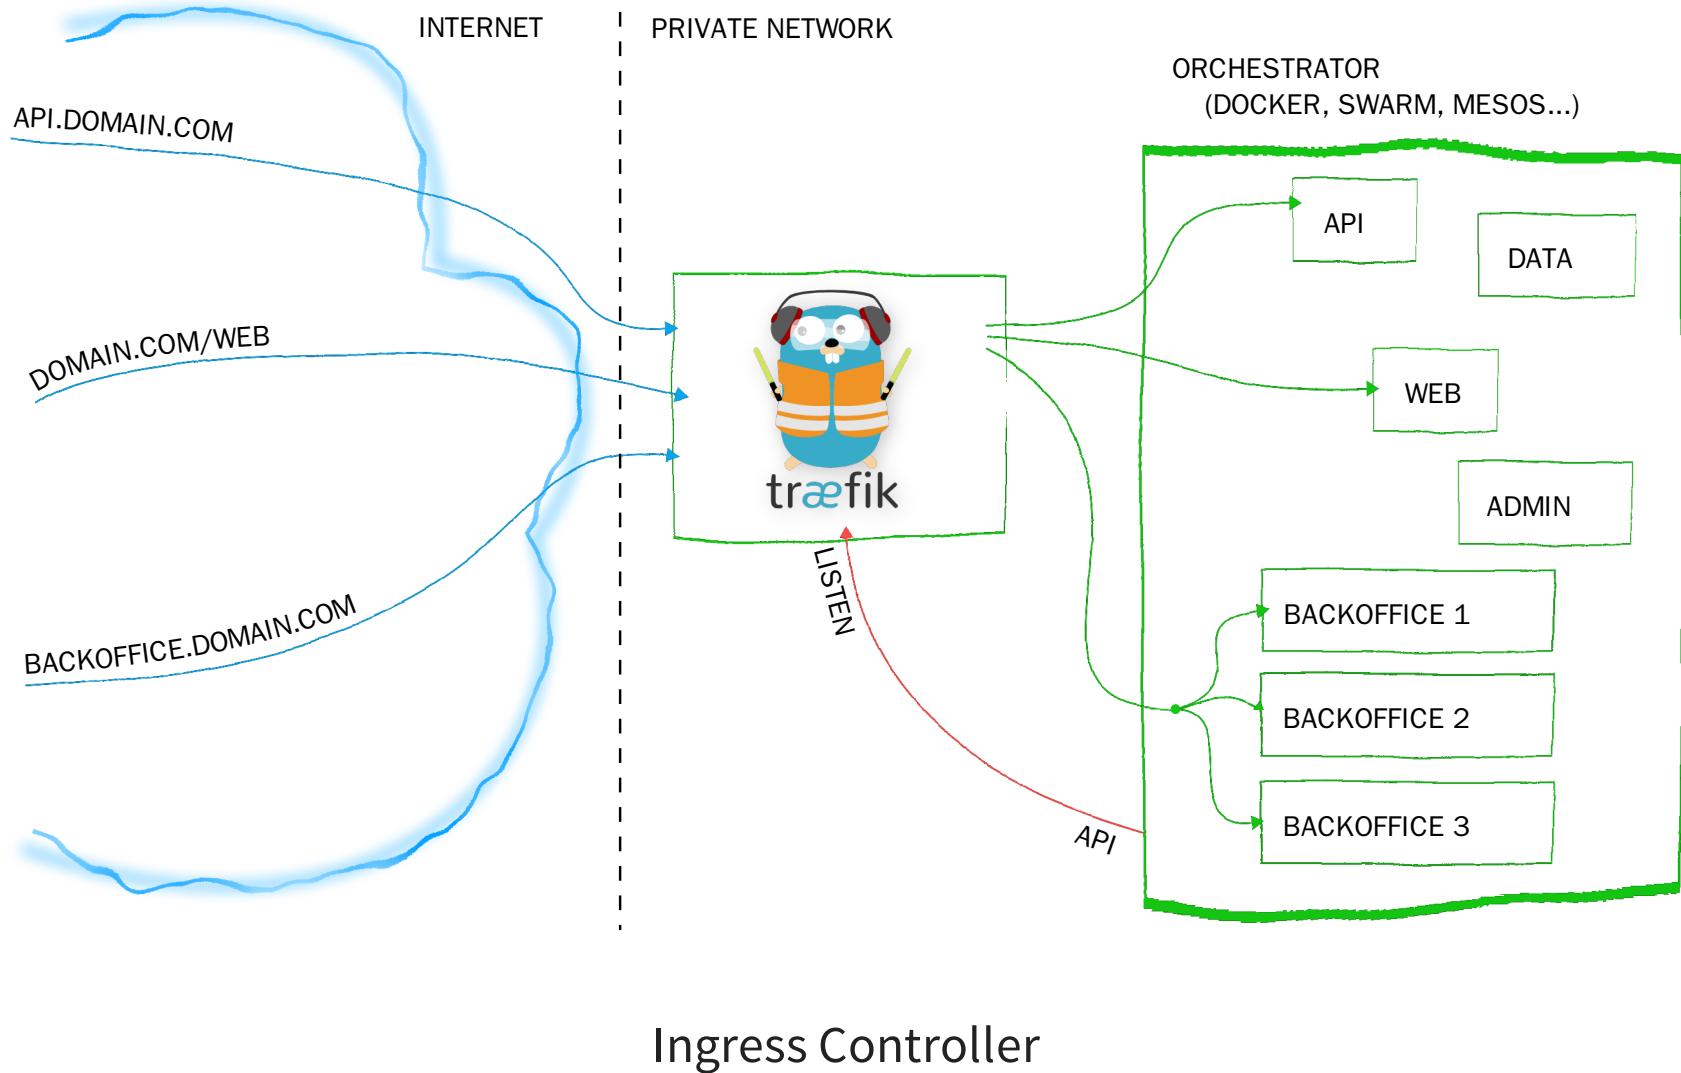
# RESISTANCE!



# RESISTANCE!



# HOW GET USER REQUESTS?



# INGRESS

Pattern

api.smacc.io/v1/users

Target App Service

users-v1

---

api.smacc.io/v2/users

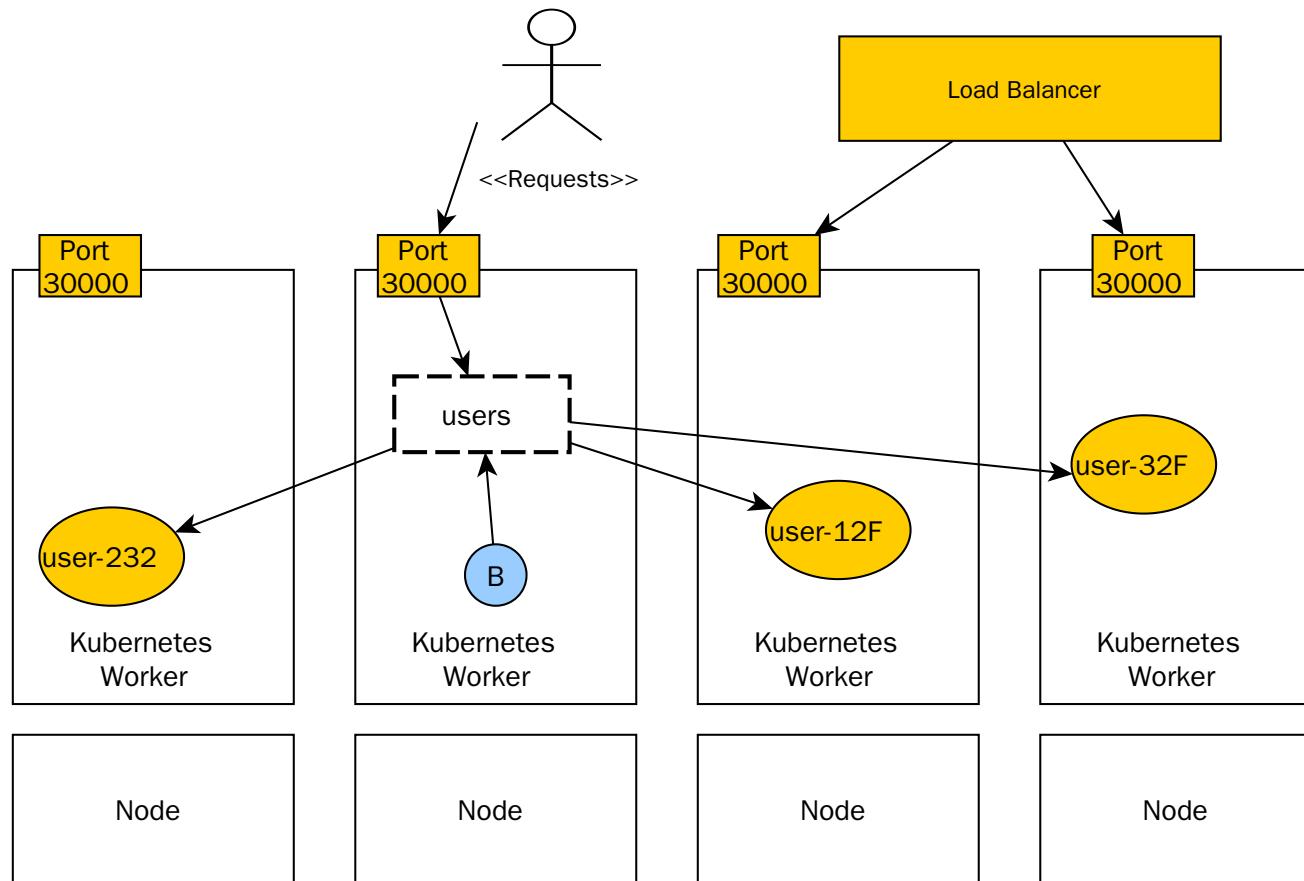
users-v2

---

smacc.io

web

# LOAD BALANCING



# SERVICE DISCOVERY

- names in DNS:

```
curl http://users/list
```

- labels:

```
name=value
```

- annotations:

```
prometheus.io/scrape: "true"
```

# SERVICE DISCOVERY

- loosely couple components
- auto-wiring with logging and monitoring

# DROP-IN

- traefik / Ingress / Envoy
- prometheus
- audit checks
- ...

# THE BEST PART

All live in git:

- all in Yaml
- integration with monitoring, alarming
- integration with ingress-controller
- ...
- Devs can forget about infrastructure... almost

DevOps Culture Dream!

# **DEPLOYMENT STRATEGIES**

# STRATEGIES

We will see:

- Replace (downtime visible)
- Rolling updates
- Blue Green
- Canary

# OTHER

We will not cover:

- Feature toggles
- A/B like
- Shadow deployment

# FIRST THE HOMEWORK

Need to support:

- liveness - am I dead?
- readiness - can I serve requests?

# KUBE LIVENESS PROBE

```
livenessProbe:  
  httpGet:  
    path: /model  
    port: 8000  
    httpHeaders:  
      - name: X-Custom-Header  
        value: Awesome  
  initialDelaySeconds: 600  
  periodSeconds: 5  
  timeoutSeconds: 18  
  successThreshold: 1  
  failureThreshold: 3
```

# LIVENESS PROBE

- our pod gets restarted
- too many restarts -> CrashLoop

# KUBE READINESS PROBE

```
readinessProbe:  
  exec:  
    command:  
    - cat  
      - /tmp/healthy  
  initialDelaySeconds: 5  
  periodSeconds: 5
```

# YOUR APP SHOULD ON STOP

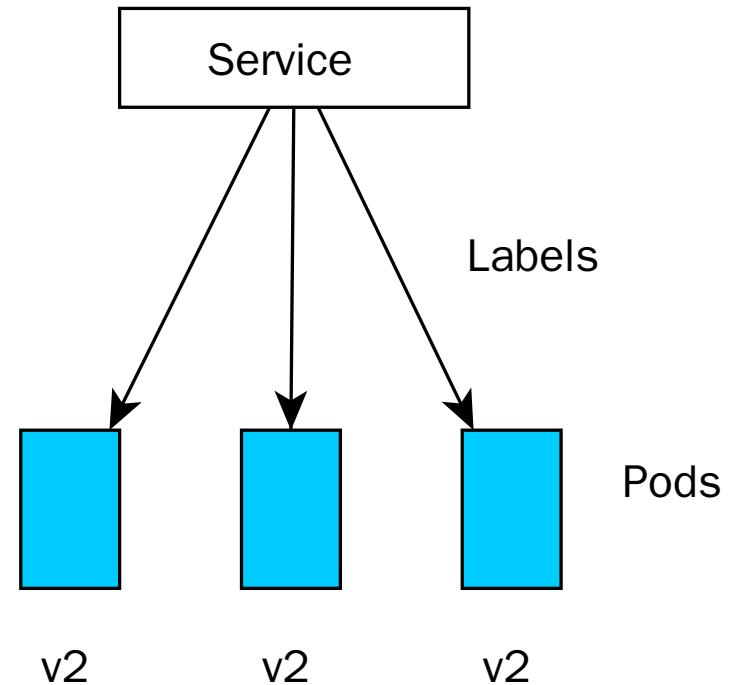
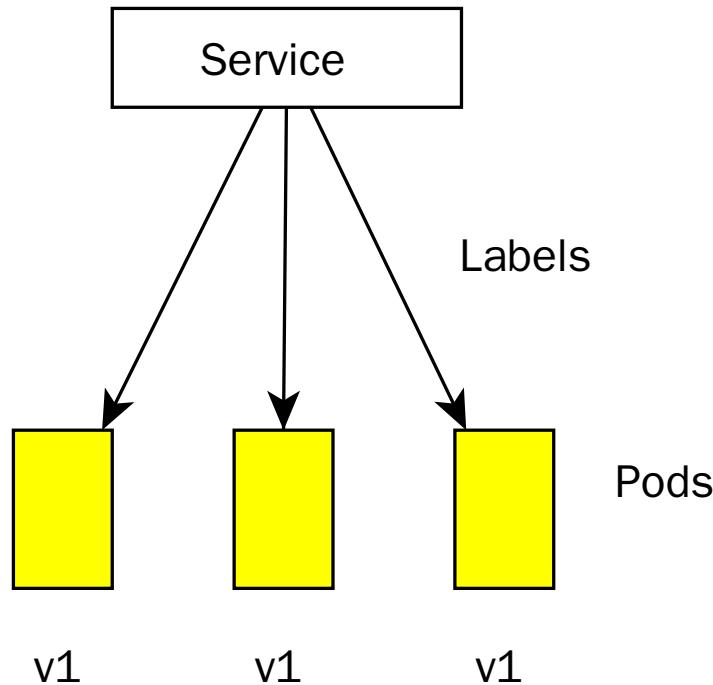
- when we get SIGTERM signal
- app gives 500 on readinessProbe
- app does not receive new requests
- app graceful shutdown
- kubernetes forces kill if 30 limit exceeded

# ALWAYS

Implement readiness for:

- ML Model-based components
- slow starting time

# DEMO - RECREATE



# DEMO - RECREATE

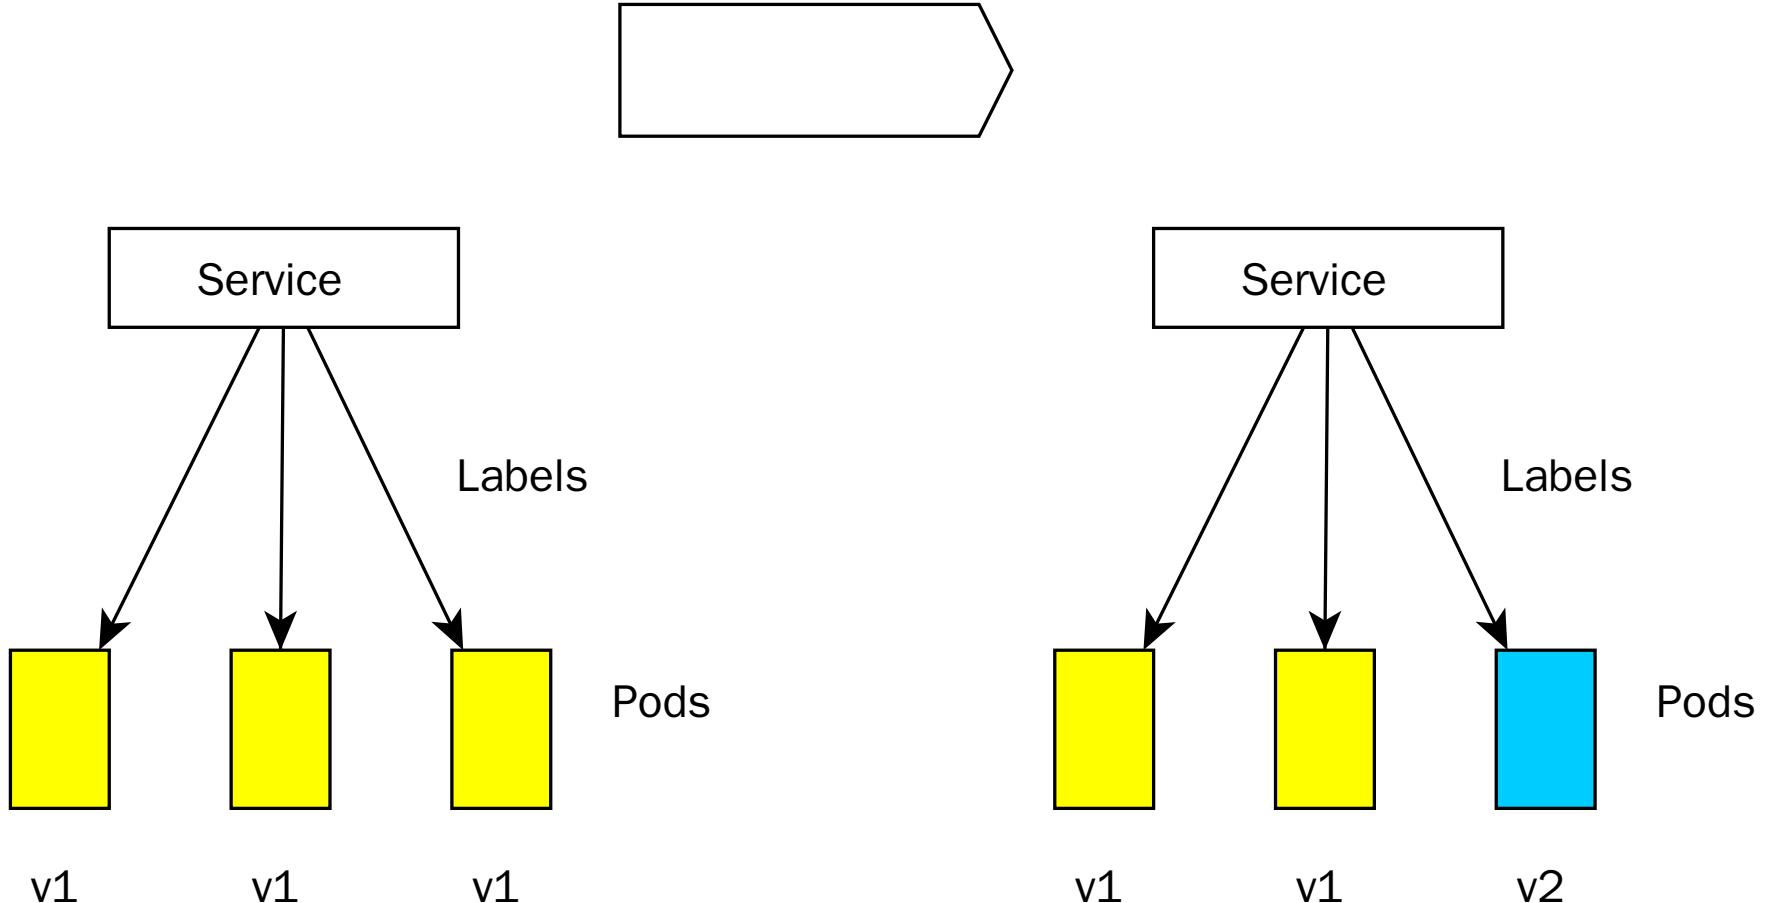
```
spec:  
  replicas: 3  
  strategy:  
    type: Recreate
```

```
kubectl set image deployment/demo-api \  
  app=wojciech11/api-status:2.0.0
```

## **DEMO - RECREATE**

- quick
- downtime visible

# DEMO - ROLLING UPDATES



# DEMO - ROLLING UPDATES

```
strategy:  
  type: RollingUpdate  
  rollingUpdate:  
    maxSurge: 2  
    maxUnavailable: 0
```

[docs](#)

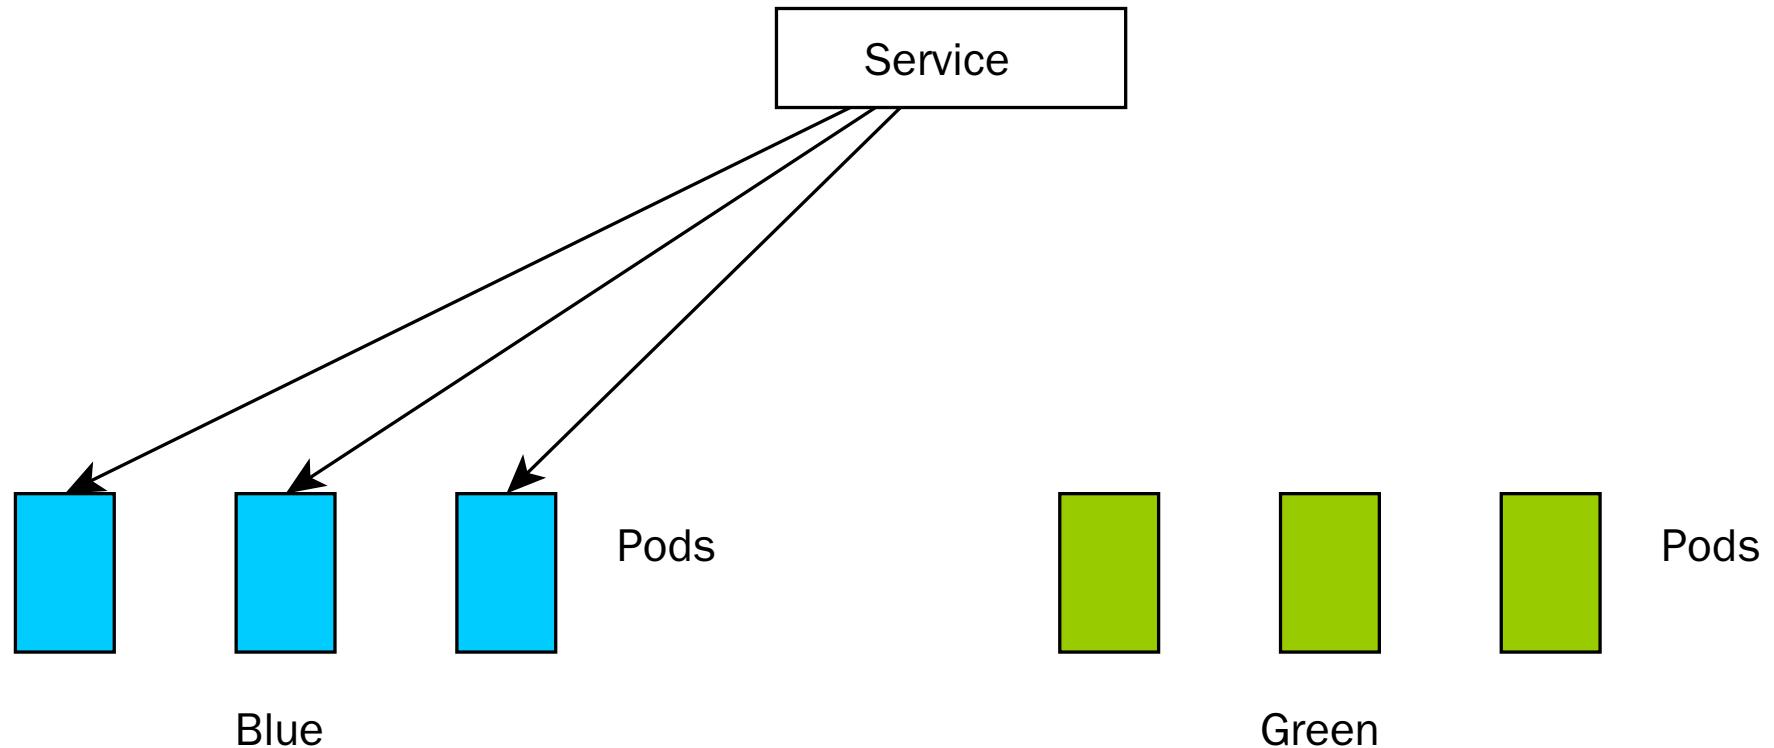
# DEMO - ROLLING UPDATES

```
kubectl set image deployment/demo-api  
  app=wojciech11/api-status:2.0.0
```

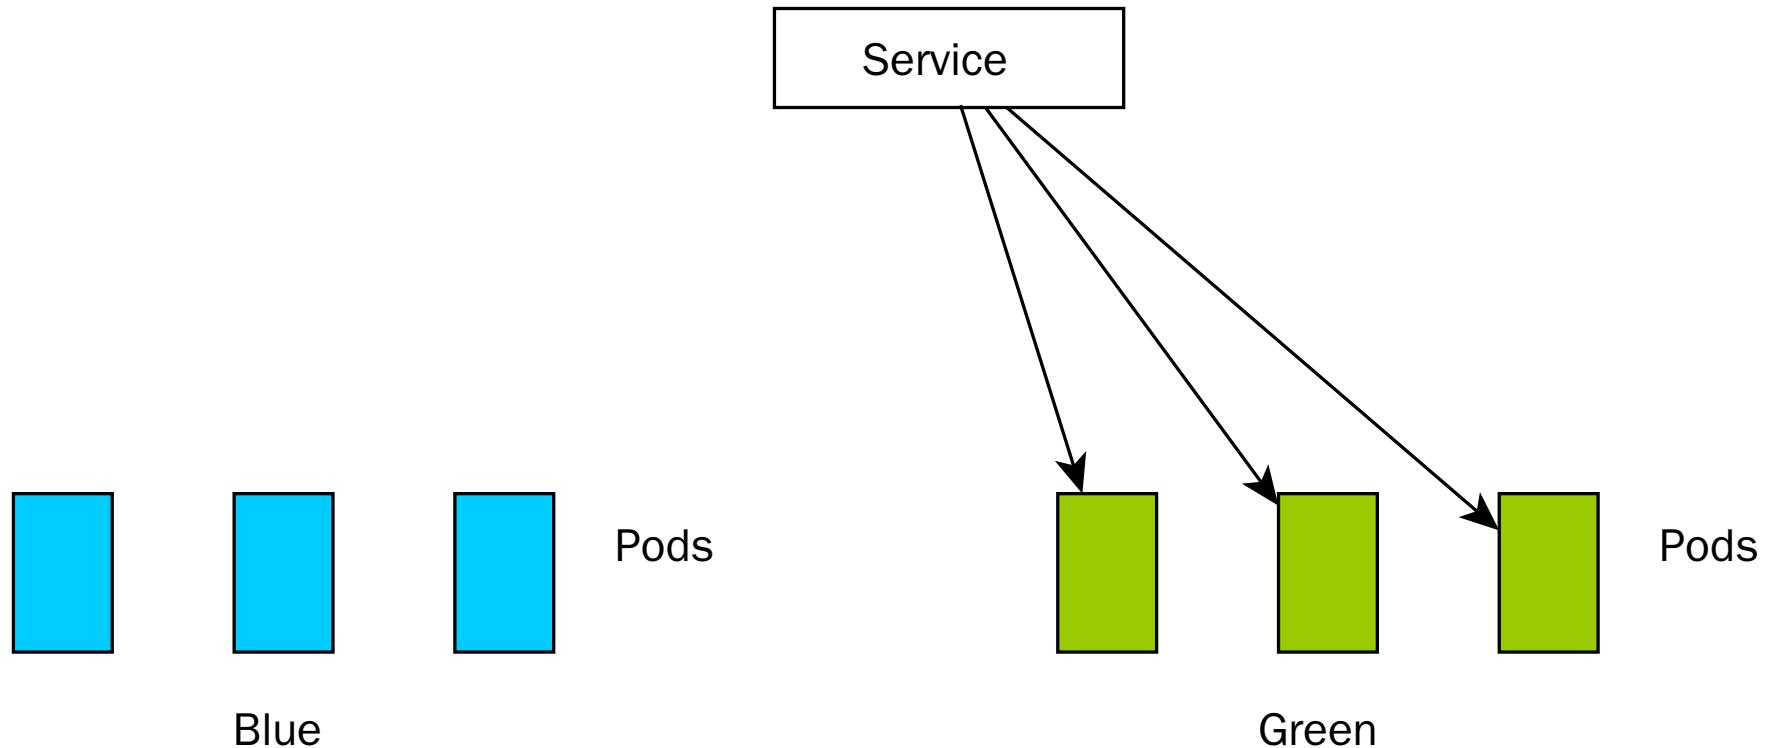
# **DEMO - ROLLING UPDATES**

- the most popular

# DEMO - GREEN/BLUE



# DEMO - GREEN/BLUE



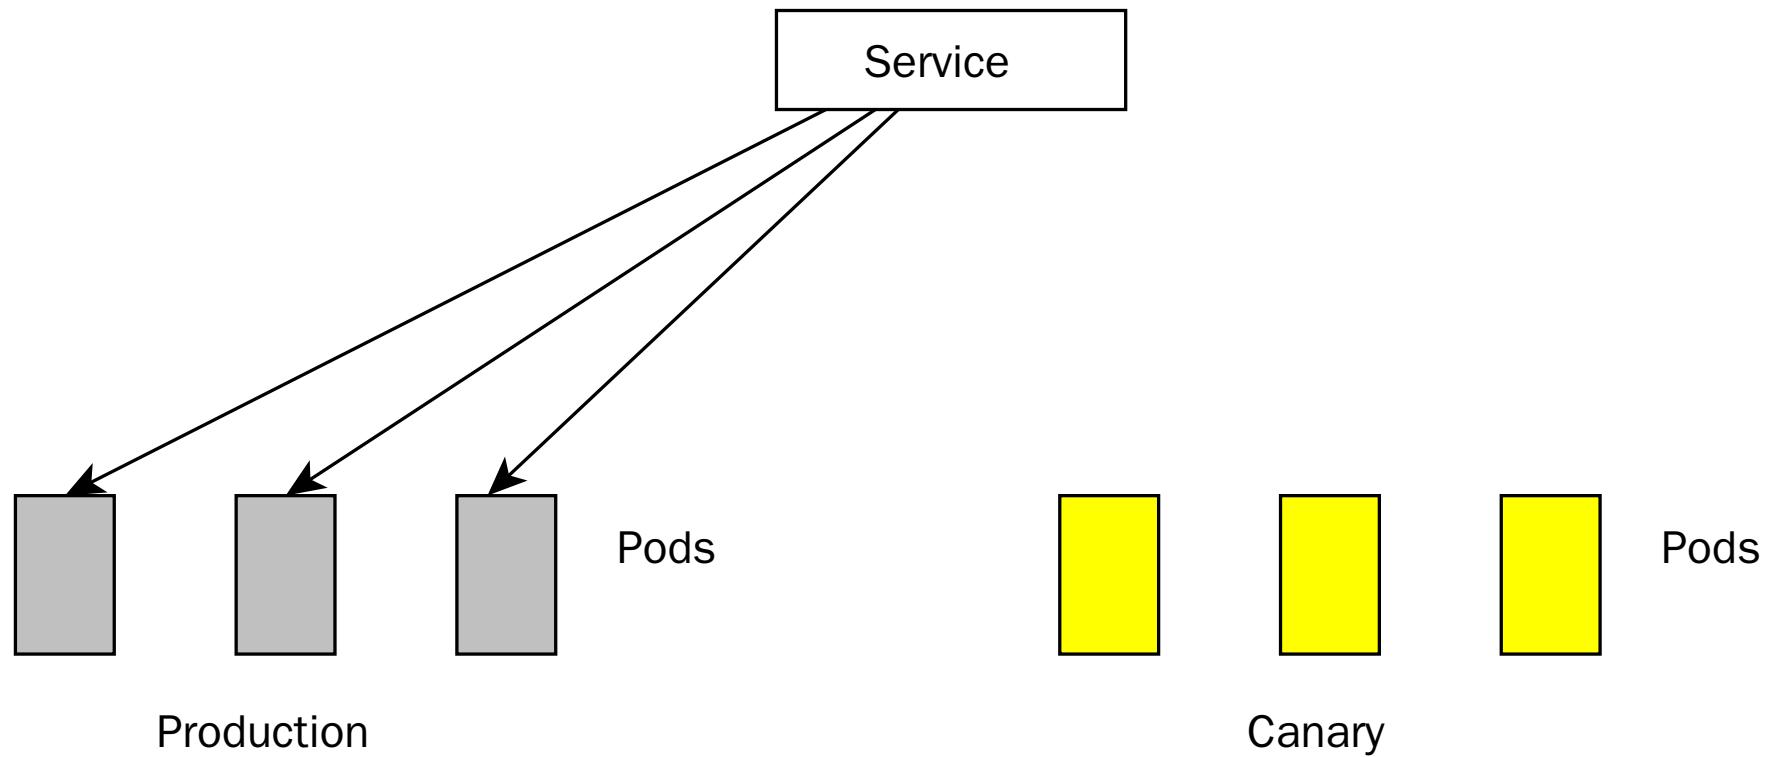
# DEMO - GREEN/BLUE

```
kubectl patch service api-status \
  -p '{"spec":{"selector":{"label": "green"}}}'
```

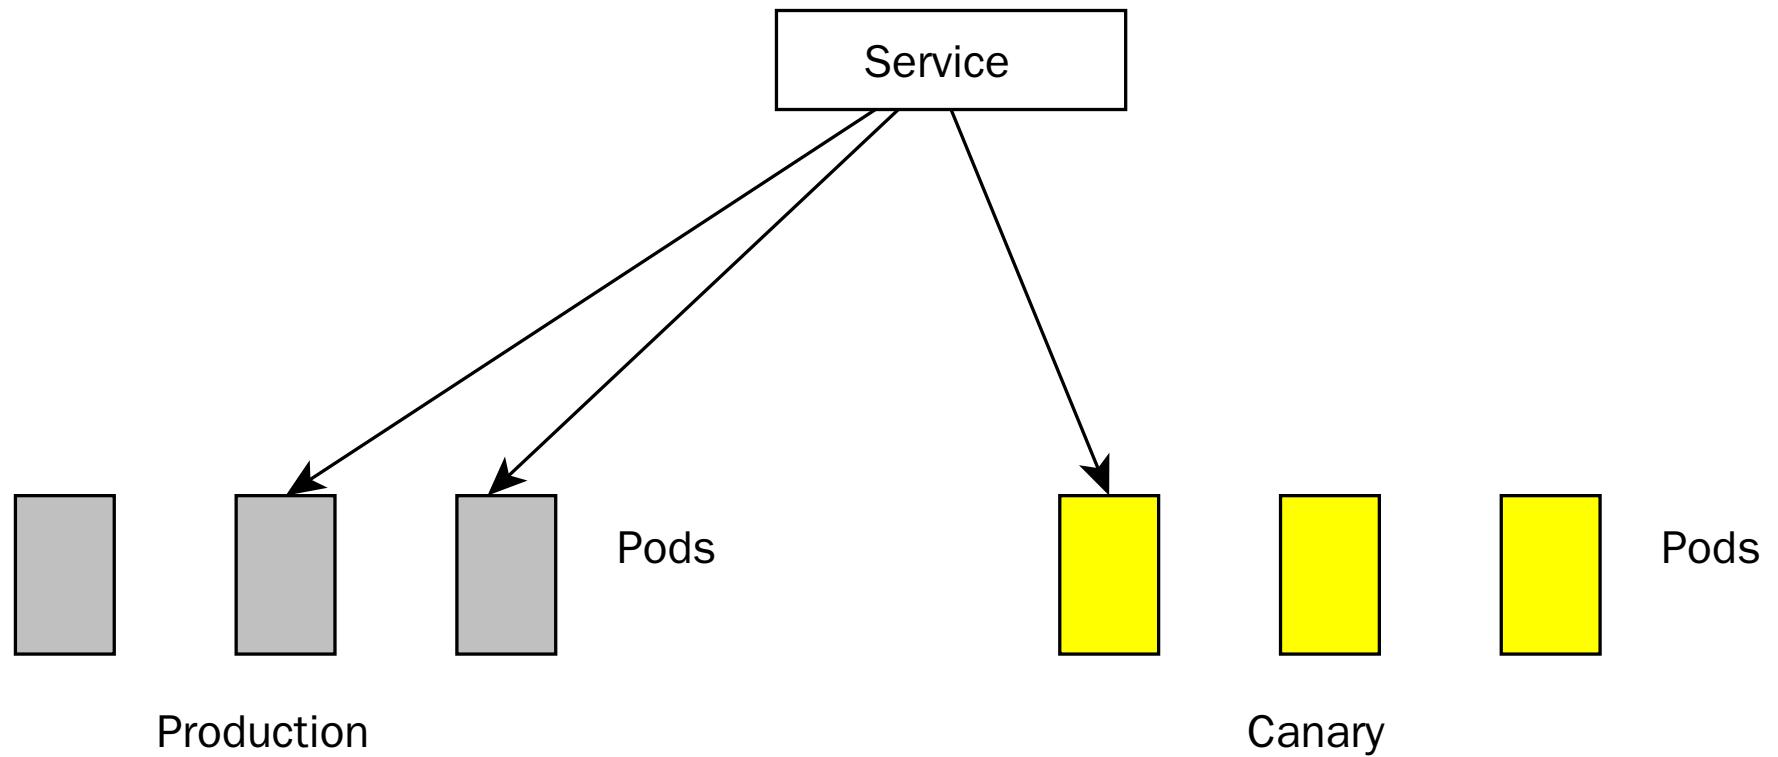
## **DEMO - GREEN/BLUE**

- For big changes
- Less common
- Might be implemented with *Ingress*

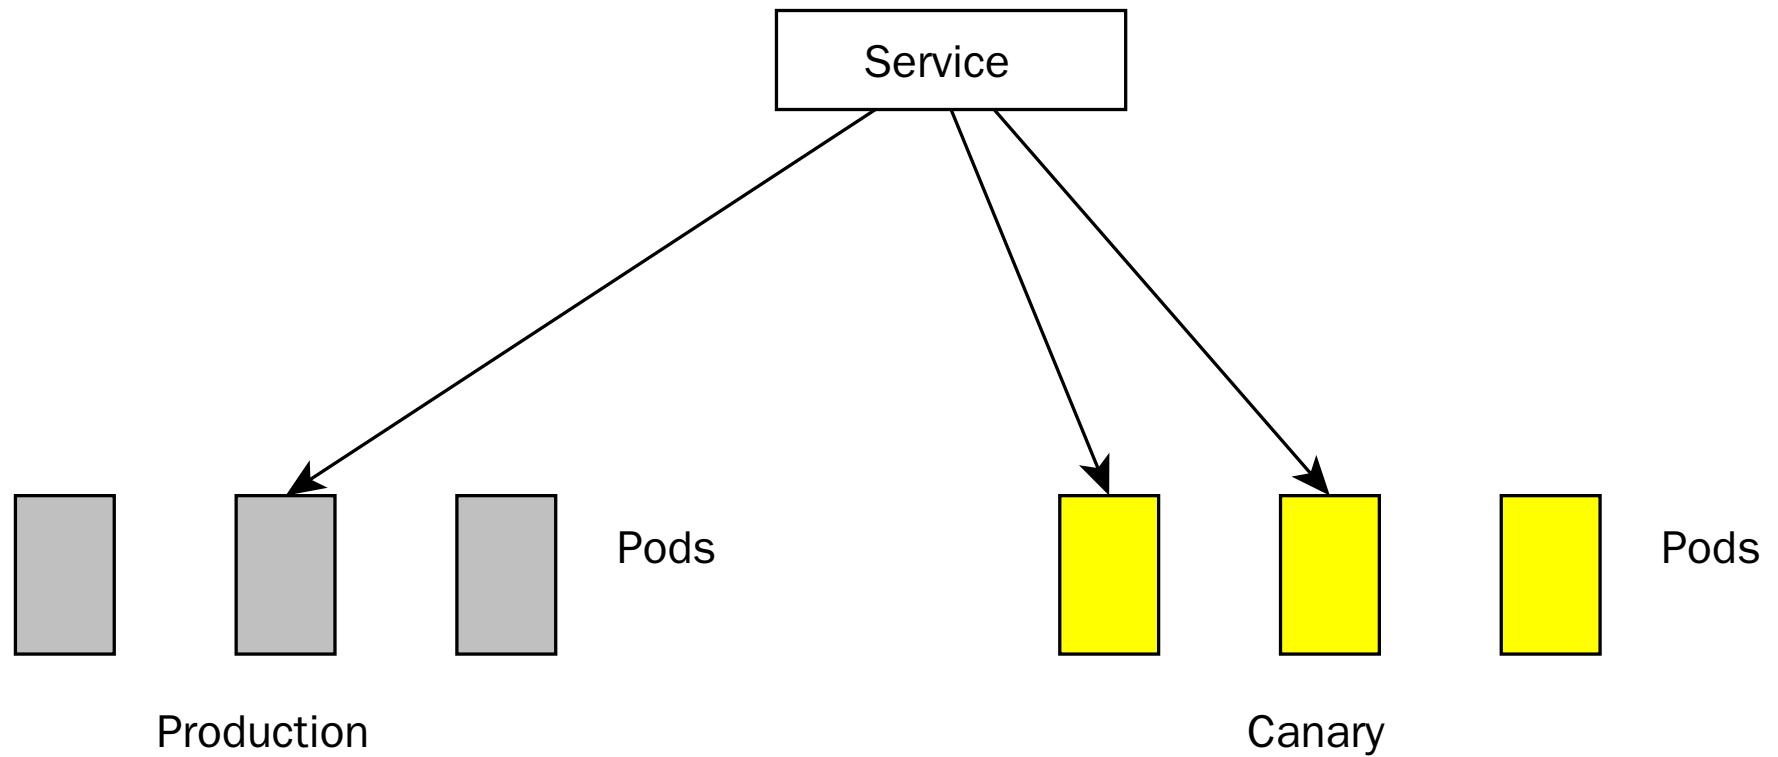
# DEMO - CANARY



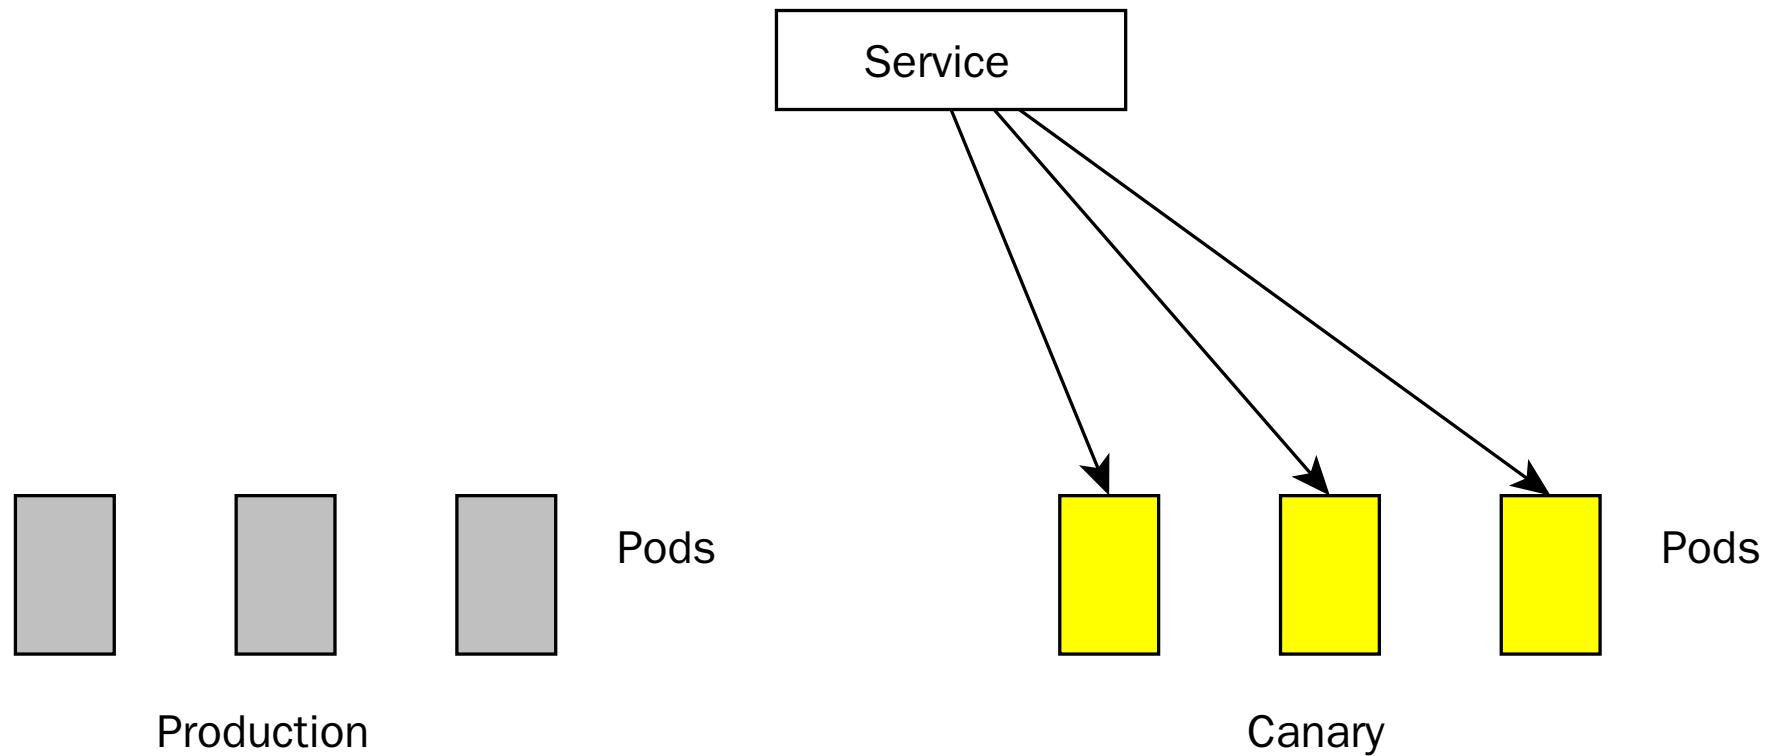
# DEMO - CANARY



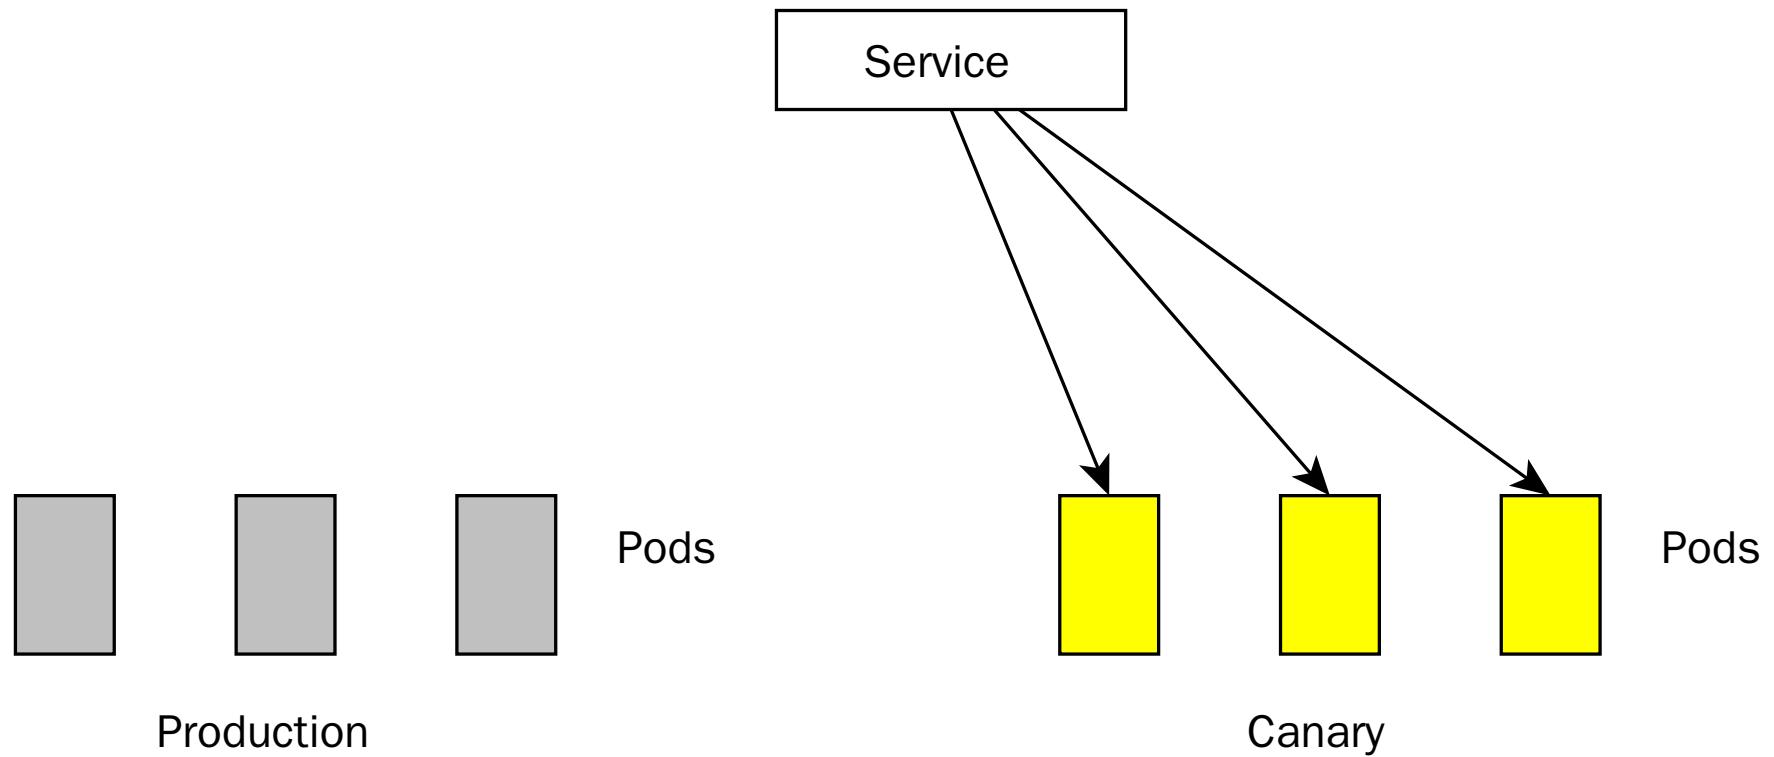
# DEMO - CANARY



# DEMO - CANARY



# DEMO - CANARY



```
kubectl scale --replicas=3 deploy/api-status-nginx-blue
kubectl scale --replicas=1 deploy/api-status-nginx-green

# no errors, let's continue
kubectl scale --replicas=2 deploy/api-status-nginx-blue
kubectl scale --replicas=2 deploy/api-status-nginx-green
```

# DEMO - CANARY

- manually
- with help of Traefik / Istio / ...

# SUMMARY

- kubernetes simple semantic
- easy deployment of your applications
- will work for any application type

# DZIĘKUJĘ. PYTANIA?

ps. We are hiring.

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# SMACC



Go



PYTORCH

TensorFlow™



amazon  
web services™



Azure



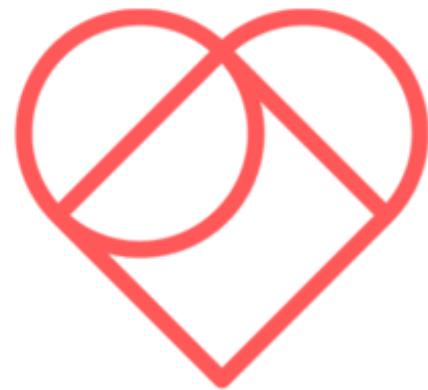
# BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



# STORY

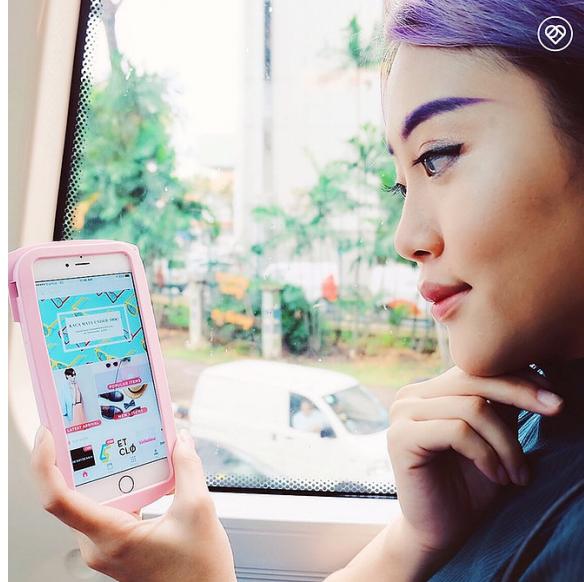
- Lyke - [12.2016 - 07.2017]
- SMACC - [10.2017 - present]



LYKE

# LYKE

- E-commerce
- Mobile-only
- 50k+ users
- 2M downloads
- Top 10 Fashion Apps  
w Google Play Store



<http://www.news.getlyke.com/single-post/2016/12/02/Introducing-the-New-Beautiful-LYKE>

Now JollyChic Indonesia

# GOOD PARTS

- Fast Growth
- A/B Testing
- Data-driven
- Product Manager,  
UI Designer,  
Mobile Dev,  
and tester - one  
body



## CHALLENGES

- 50+ VMs in Amazon, 1 VM - 1 App, idle machine
- Puppet, hilarious (manual) deployment process
- Fear
- Forgotten components
- sometimes performance issues

# SMACC

---

## Hypatos

# SMACC

- Machine Learning FinTech
- SaaS and API platform
- From Enterprise (Deutsche Bank, AoK) to SME
- Well-known FinTech Startup in Germany



# STORY

- Legacy on AWS, experiments with AWS ECS :/
- Self-hosted K8S on ProfitBricks
- Get to Microsoft ScaleUp, welcome Azure
- Luckily - Azure-Kubernetes-Service

# DIFFERENCE

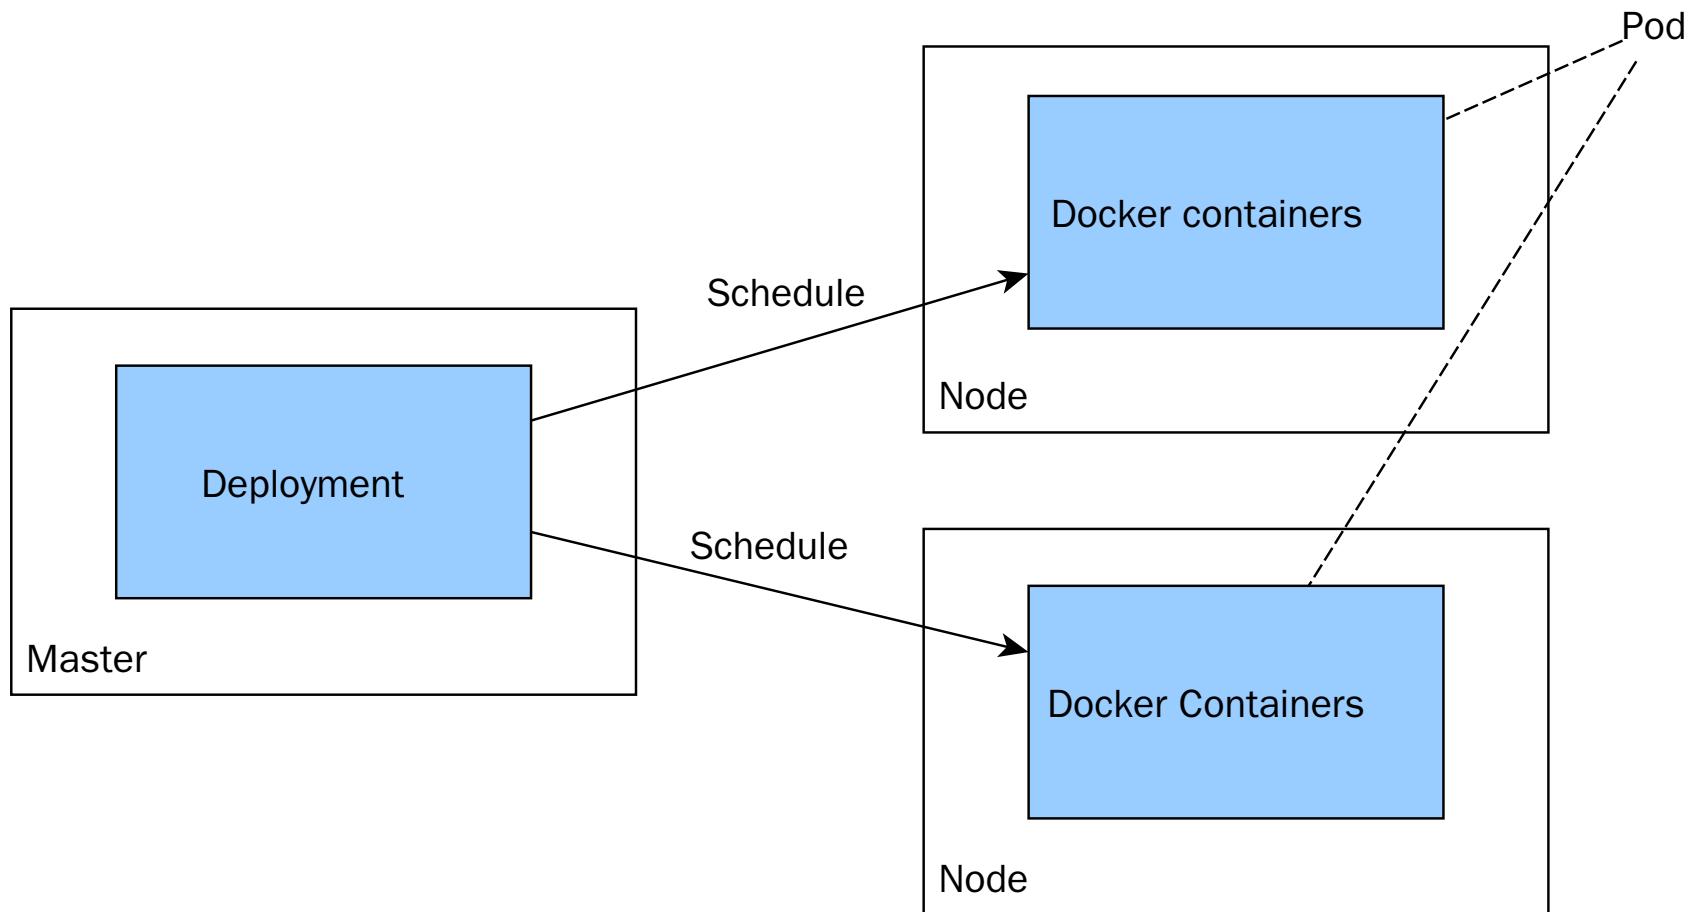


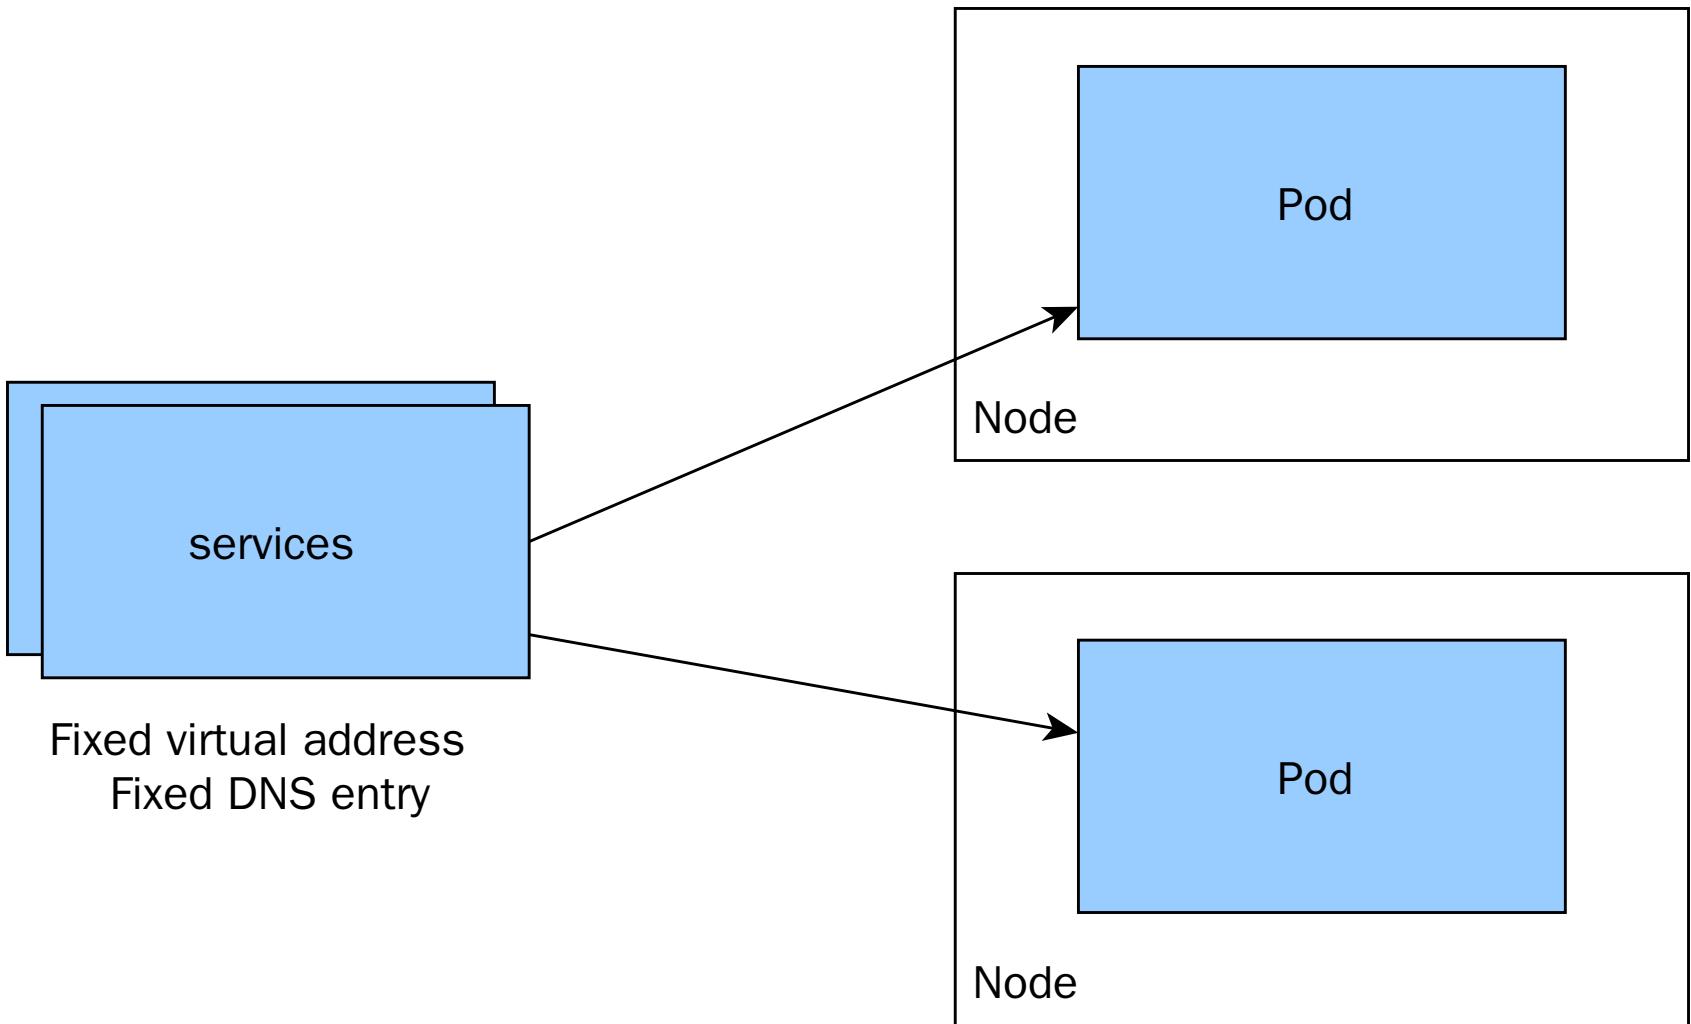
- Two teams in Berlin and Warsaw
- Me in Warsaw

## APPROACH

- Simplify, Simplify
- Hide K8S magic
- git tag driven Continuous Deployment

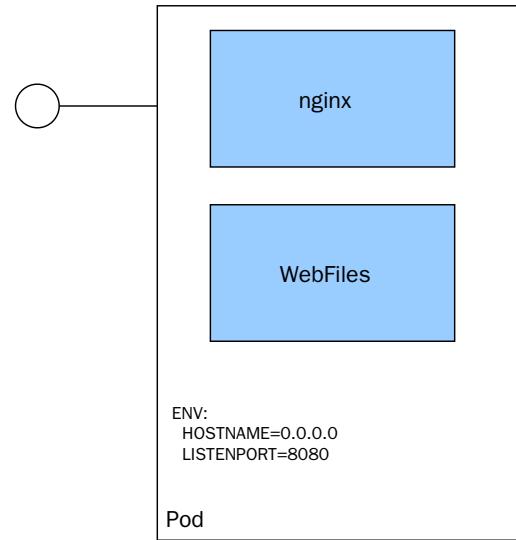
# KUBERNETES CONCEPTS



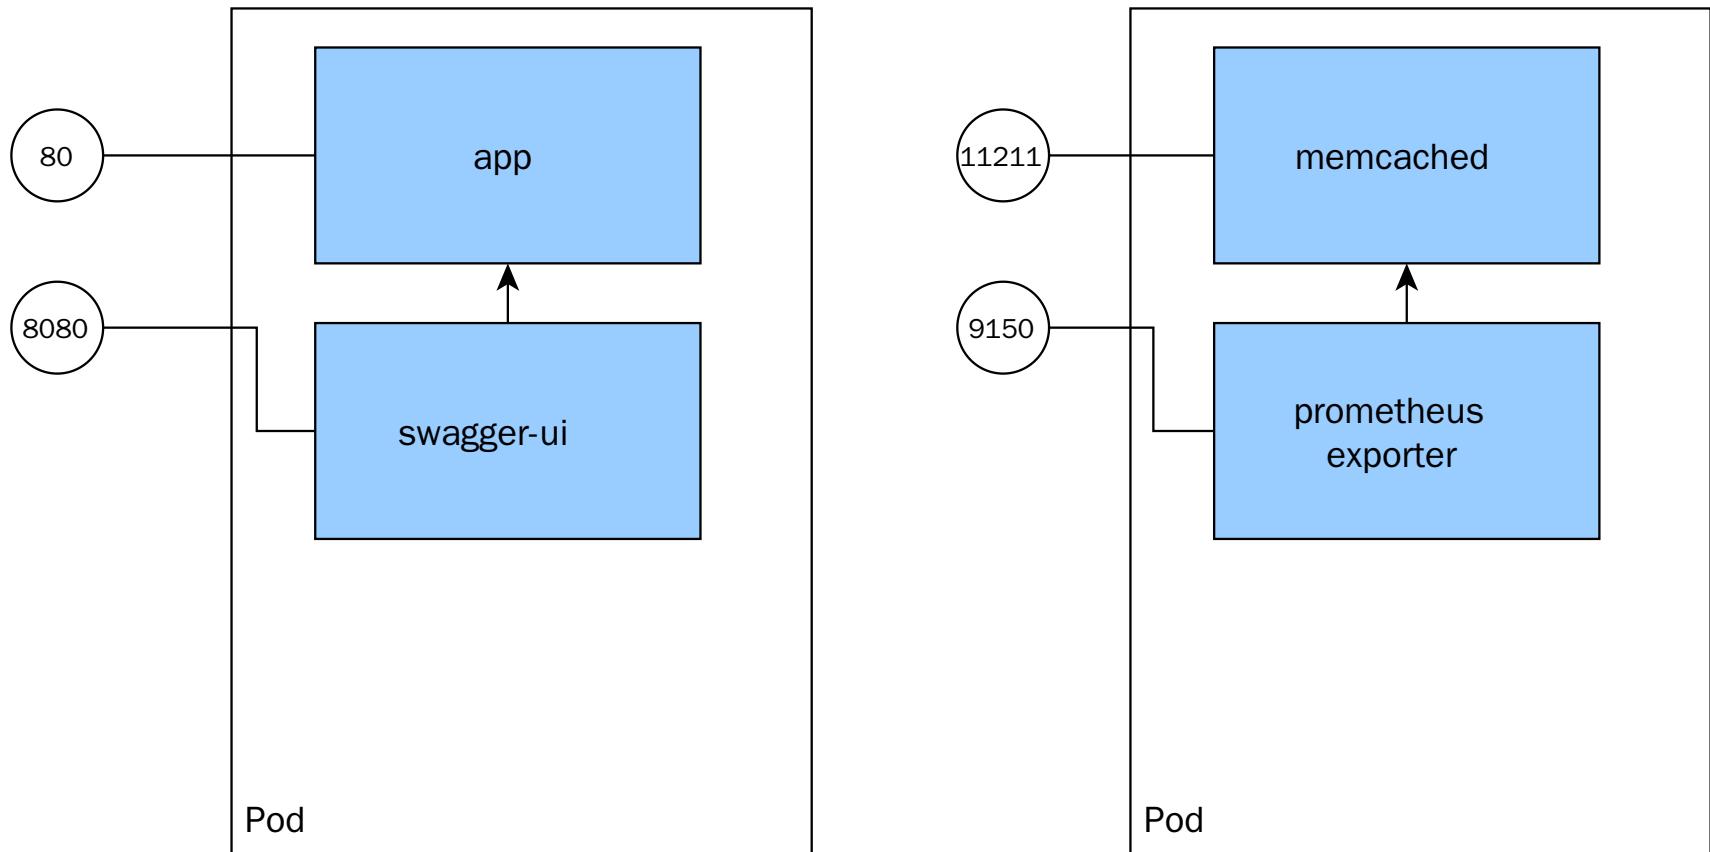


# PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports



# SIDE-CARS



# BASIC CONCEPTS

Name	Purpose	
Service	Interface	Entry point (Service Name)
Deployment	Factory	How many pods, which pods
Pod	Implementation	1+ docker running